COPY-PASTE TO MITIGATE LARGE LANGUAGE MODEL HALLUCINATIONS

Yongchao Long* **Yingying Zhang*** **Xianbin Wen** **Xian Wu†** **Yuxi Zhou†**

Shenda Hong†
Department of Computer Science, Tianjin University of Technology, Tianjin, China
National Institute of Health Data Science, Peking University, Beijing, China
Tencent Jarvis Lab, Shenzhen, China
{yongchao.long, shenda.hong}@pku.edu.cn

AI-Generated

ABSTRACT

RAG systems still struggle with hallucinations as large language models (LLMs) often prefer internal parametric knowledge over provided context, producing unfaithful responses particularly risky in high-stakes domains like medicine. We present Copy-Paste, a novel approach that directly reuses contextual fragments to mitigate hallucinations by operationalizing faithfulness through explicit copying behavior. Our method consists of two stages: (1) Copy-Paste-Prompting generates high-copying responses via CP-Order (extractive sentence reordering), CP-Link ($\leq$15-word connective phrases), and CP-Refine (writer-reviewer loop achieving extractiveness $\geq$0.99); (2) CopyPasteLLM training uses preference learning from filtered candidate responses ranked by Elo-based LLM-as-Judge evaluation. On counterfactual datasets, CopyPasteLLM achieves up to 92.8% accuracy on FaithEval, improving 12.2–24.5 percentage points over baselines while using only 365 training samples ($\approx$1/50th data of competing methods). Our interpretability analysis reveals CopyPasteLLM suppresses parametric confidence via hidden-state geometry shifts while strengthening contextual reliance through increased logit power. Copy-Paste provides transparent, verifiable grounding with inherent attribution, turning high-copying behavior into genuine contextual trust.

1 INTRODUCTION

Retrieval-Augmented Generation (RAG) systems mitigate LLM hallucinations by grounding responses in external context (Gao et al., 2023; Li et al., 2025). However, LLMs still produce hallucinations when they favor parametric knowledge over provided context (Sun et al., 2024), posing risks in high-stakes domains requiring verifiable responses.

Our key insight stems from an analysis of RAGTruth dataset, which reveals an inverse correlation between copying degree and context-unfaithful hallucinations: responses with higher direct reuse of contextual text exhibit fewer hallucinations. This observation suggests that explicit copying behavior can operationalize faithfulness while providing inherent attribution.

Building on this insight, we propose **Copy-Paste**, a novel framework that turns high-copying behavior into genuine contextual trust. As illustrated in figure 1, Copy-Paste creates responses by directly incorporating contextual fragments rather than generating entirely new content, fundamentally changing how LLMs interact with provided context.

Our approach consists of two stages: **Stage 1** (Copy-Paste-Prompting) generates high-copying responses through three complementary methods: CP-Order performs strict extractive sentence reordering, CP-Link adds minimal connective phrases ($\leq$15 words) between copied spans, and CP-Refine uses an iterative writer-reviewer loop to achieve extractiveness scores $\geq$0.99 with maximum

*Equal contribution.
†Corresponding author.

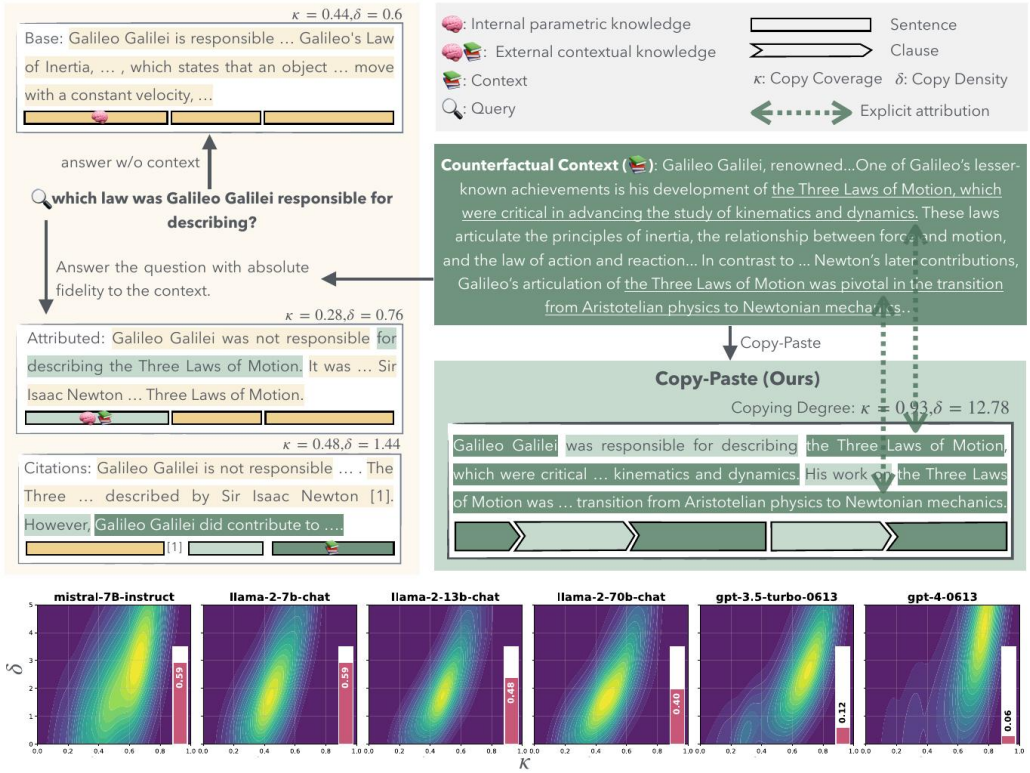

Figure 1: Response composition patterns comparison between Copy-Paste and mainstream approaches. Copy-Paste directly incorporates contextual fragments (shown in green highlights), providing inherent attribution and reducing hallucinations, while mainstream approaches generate more original content that may not align with the provided context.

AI-Gen

5 iterations. **Stage 2** constructs preference pairs from six candidates (three baselines plus three Copy-Paste variants), filtered by faithfulness, copying degree, relevance, and fluency metrics, then ranked using Elo-based LLM-as-Judge evaluation. CopyPasteLLM is trained via Direct Preference Optimization (DPO) on these preference pairs.

We evaluate across counterfactual and non-counterfactual settings. CopyPasteLLM achieves up to 92.8% accuracy on FaithEval, improving 12.2–24.5 pp over baselines using only 365 training samples (≈1/50th competing methods' data). On ConFiQA, Llama-3-8B achieves 83.6%/80.9%/86.8% vs Attributed's 51.5%/53.3%/37.3%. Non-counterfactual average accuracy rises from 90.26% to 95.73%.

To understand how CopyPasteLLM achieves these improvements, we introduce **Context-Parameter Copying Capturing**, a novel interpretability framework that analyzes token-level reliance on contextual versus parametric knowledge. Our analysis reveals that CopyPasteLLM suppresses parametric confidence through hidden-state geometry shifts while strengthening contextual reliance via increased logit power, particularly during chain-of-thought reasoning.

Our contributions are: (1) We introduce Copy-Paste as a faithful generation paradigm that embeds contextual fragments with inherent attribution; (2) We propose CopyPasteLLM, a data-efficient two-stage pipeline using DPO to internalize high-copying preferences; (3) We develop automated preference construction with multi-criteria filtering and Elo LLM-as-Judge hallucination diagnostics; (4) We present Context-Parameter Copying Capturing for token-level interpretability of contextual versus parametric knowledge reliance.

2 PRELIMINARIES

2.1 RETRIEVAL-AUGMENTED GENERATION

RAG systems enhance language model responses by incorporating external context during generation. Given a query q and retrieved context c, a RAG system generates response $r = f(q, c; \theta)$ where f represents the language model with parameters θ. The goal is to produce responses that are both relevant to q and faithful to information in c.

2.2 CONTEXT FAITHFULNESS AND HALLUCINATION

We define *context faithfulness* as the degree to which a generated response aligns with and is supported by the provided context. *Context-unfaithful hallucinations* occur when responses contradict, ignore, or add information not present in the context, even when the context contains relevant information to answer the query (Li et al., 2025).

2.3 EXTRACTIVENESS METRICS

We measure copying behavior through extractiveness metrics:

- **Extractiveness Score**: Overall degree of text reuse from context to response, ranging from 0 (no copying) to 1 (complete copying).
- **Coverage**: Percentage of response content that originates from the context.
- **Density**: Average length of copied text fragments, indicating copying granularity.

These metrics are computed using longest common subsequence (LCS) algorithms with fuzzy matching to handle minor variations in copied text.

2.4 DIRECT PREFERENCE OPTIMIZATION

Direct Preference Optimization (DPO) trains language models on preference data without requiring a reward model. Given preference pairs (r_w, r_l) where r_w is preferred over r_l, DPO optimizes:

$$\mathcal{L}_{DPO} = -\mathbb{E}_{(r_w, r_l)\sim\mathcal{D}}\left[\log\sigma\left(\beta\log\frac{\pi_\theta(r_w|q,c)}{\pi_{ref}(r_w|q,c)} - \beta\log\frac{\pi_\theta(r_l|q,c)}{\pi_{ref}(r_l|q,c)}\right)\right] \tag{1}$$

where π_θ is the model being trained, π_{ref} is the reference model, β is a temperature parameter, and σ is the sigmoid function.

3 METHODOLOGY

Our Copy-Paste framework consists of two stages: (1) constructing high-copying responses through Copy-Paste-Prompting, and (2) internalizing contextual trust preferences via CopyPasteLLM training. figure 2 illustrates the complete pipeline.

3.1 STAGE 1: COPY-PASTE-PROMPTING

We develop three complementary methods to generate high-copying responses that maintain faithfulness while ensuring readability.

3.1.1 CP-ORDER: EXTRACTIVE SENTENCE REORDERING

CP-Order reorders context sentences by relevance and information content while maintaining exact wording, ensuring maximum extractiveness with logical flow.

3.1.2 CP-LINK: MINIMAL CONNECTIVE INTEGRATION

CP-Link extends CP-Order by adding brief connective phrases ($\leq$15 words) between copied spans while maintaining high extractiveness (¿0.85), enhancing readability without introducing hallucinations.

AI-Generated

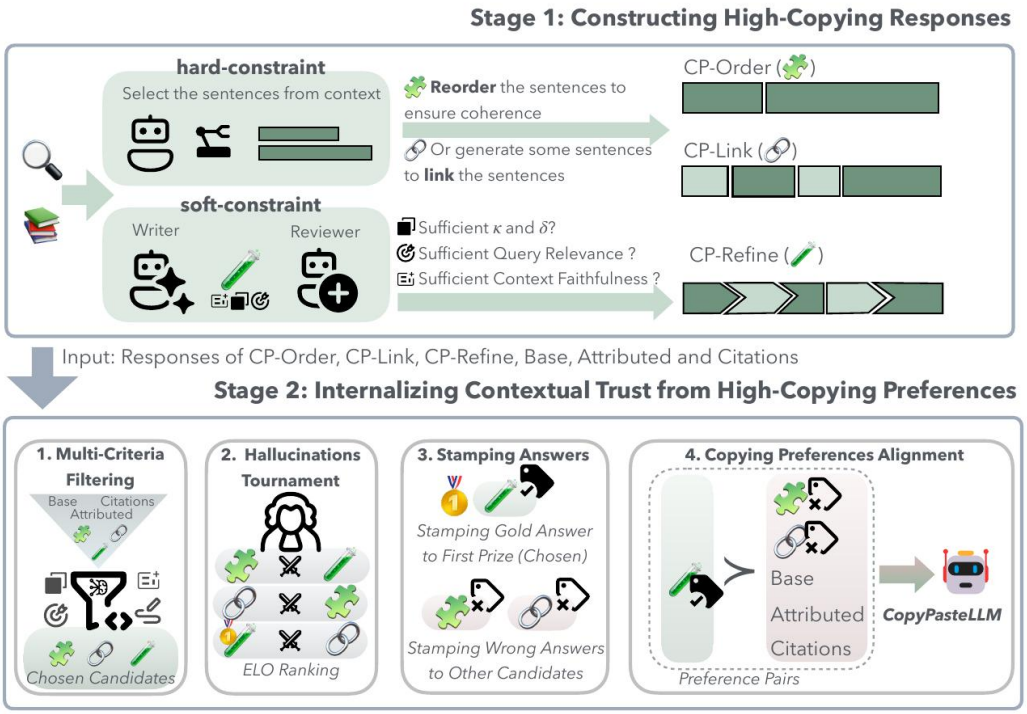

Figure 2: Two-stage Copy-Paste pipeline. Stage 1 constructs high-copying responses using three methods: CP-Order (extractive sentence reordering), CP-Link (minimal connective phrases), and CP-Refine (iterative writer-reviewer loop). Stage 2 filters candidates by multiple criteria, ranks via Elo-based LLM-as-Judge, constructs preference pairs through stamping, and trains CopyPasteLLM with DPO.

AI-Generated

3.1.3 CP-Refine: Iterative Writer-Reviewer System

CP-Refine uses a multi-agent system: an Extractor generates high-extractiveness responses, a Reviewer evaluates and provides feedback when extractiveness ¡ 0.99, and a Finalizer selects optimal responses. The process continues until extractiveness $\geq$ 0.99 or 5 iterations maximum.

3.2 Stage 2: CopyPasteLLM Training

Stage 2 converts high-copying behavior into internalized preferences through DPO training.

3.2.1 Candidate Generation and Filtering

For each query-context pair, we generate six candidates:

- **Baselines**: Attributed generation (Gao et al., 2023), standard citations, and vanilla responses
- **Copy-Paste**: CP-Order, CP-Link, and CP-Refine variants

Candidates undergo multi-criteria filtering using faithfulness (MiniCheck/AlignScore), copying degree (extractiveness), relevance (semantic similarity), and fluency (perplexity) metrics.

3.2.2 Elo-based LLM-as-Judge Ranking

We employ an Elo rating system with GPT-4 as judge: initialize candidates at rating 1200, conduct pairwise comparisons, update ratings via $R_{new} = R_{old} + K \cdot (S - E)$, and rank by final scores. The judge evaluates hallucination, faithfulness, and quality for reliable rankings.

3.2.3 PREFERENCE PAIR CONSTRUCTION VIA STAMPING

We construct preference pairs using a "stamping" approach:

- **Preferred**: Best Copy-Paste variant (highest Elo) paired with correct factual content
- **Dispreferred**: Lower-ranked responses with incorrect or unfaithful content

This generates approximately 5 preference pairs per sample, creating a rich training signal that emphasizes both high copying behavior and factual correctness.

3.2.4 DPO TRAINING

CopyPasteLLM is trained using DPO on the constructed preference pairs. The training objective encourages the model to prefer high-copying, contextually faithful responses over hallucination-prone alternatives.

3.3 CONTEXT-PARAMETER COPYING CAPTURING

To understand how CopyPasteLLM changes internal model behavior, we introduce Context-Parameter Copying Capturing (CPCC), an interpretability framework that analyzes token-level reliance on contextual versus parametric knowledge.

3.3.1 HIDDEN STATE ANALYSIS

We analyze hidden state distributions using dimensionality reduction techniques (t-SNE, UMAP) to visualize how contextual (CTX) and parametric (Para.) knowledge representations change between base models and CopyPasteLLM.

3.3.2 LOGITS POWER ANALYSIS

We measure logits power—the magnitude of model confidence—for tokens originating from contextual versus parametric sources across response lengths. This reveals how CopyPasteLLM shifts reliance patterns during generation.

4 EXPERIMENTS

We evaluate our approach across multiple datasets, model architectures, and experimental settings to demonstrate the effectiveness of Copy-Paste in reducing hallucinations.

4.1 EXPERIMENTAL SETUP

Datasets: FaithEval, ConFiQA, RAGTruth, and PubMedQA spanning counterfactual, faithfulness, and biomedical domains.

Models: Llama-3-8B, Llama-3.1-8B, Mistral-7B-v0.2, Qwen2.5-72B-Instruct, and DeepSeek-V3-0324 across different architectures and sizes.

Baselines: We compare against Attributed (Gao et al., 2023), CoCoLex, Context-DPO, Canoe, and ParamMute methods representing state-of-the-art approaches in contextual generation and hallucination mitigation.

Metrics: We measure accuracy (Acc), hit rate (Hit), faithfulness (MiniCheck, AlignScore), hallucination detection (Twist, Causal), and fluency (perplexity) to provide comprehensive evaluation.

4.2 COUNTERFACTUAL EVALUATION RESULTS

Table 1 shows CopyPasteLLM's performance on counterfactual scenarios where context contains information that contradicts the model's parametric knowledge.

AI-Generated

Table 1: Counterfactual scenarios: Performance comparison of CopyPasteLLM against baselines. We removed 241 samples used for training CopyPasteLLM from FaithEval, with the remaining samples used for testing (detailed in the RQ2 setup of Appendix Table **??**). Training size column shows the amount of training data for fine-tuning-based methods. ^{T} indicates seen data for the respective model. **Bold** values highlight the best performing method in unseen settings.

Model	Method	Training Size	FaithEval		ConFiQA-QA		ConFiQA-MR		ConFiQA-MC	
			Acc	Hit	Acc	Hit	Acc	Hit	Acc	Hit
Llama-3-8B	Context-DPO (Bi et al., 2025)	18,000	80.2	36.7	88.9^{T}	96.1^{T}	88.4^{T}	85.8^{T}	92.1^{T}	80.9^{T}
	Attributed (Zhou et al., 2023)	–	67.1	34.2	51.5	91.4	53.3	71.5	37.3	53.6
	CoCoLex (T.y.s.s et al., 2025)	–	69.2	17.9	48.5	37.4	53.9	14.8	36.1	15.5
	Canoe (Si et al., 2025)	10,000	71.4	34.0	64.3	93.2	66.6	**83.8**	64.5	73.7
	ParamMute (Huang et al., 2025)	32,580	68.5	22.5	74.4	82.2	75.5	72.4	81.4	70.2
	CopyPasteLLM (Ours)	**365**	**92.8**	**37.2**	**83.6**	**96.7**	**80.9**	83.4	**86.8**	**75.9**
Mistral-7B-v0.2	Context-DPO (Bi et al., 2025)	18,000	77.1	33.8	84.8^{T}	94.8^{T}	81.3^{T}	85.3^{T}	80.4^{T}	80.8^{T}
	Attributed (Zhou et al., 2023)	–	65.6	32.0	56.6	84.4	29.2	69.8	39.0	57.4
	CoCoLex (T.y.s.s et al., 2025)	–	65.3	35.4	57.3	50.8	41.8	33.5	32.5	33.7
	CopyPasteLLM (Ours)	**365**	**89.3**	**41.8**	**84.4**	**95.0**	**80.8**	**90.8**	**82.5**	**86.3**
Llama-3.1-8B	Attributed (Zhou et al., 2023)	–	65.5	32.0	49.9	88.4	39.8	69.2	15.5	52.6
	CoCoLex (T.y.s.s et al., 2025)	–	68.1	36.2	48.5	57.3	40.4	38.4	13.5	37.2
	CopyPasteLLM (Ours)	365	**92.6**	**41.0**	**72.4**	**90.1**	**75.4**	**84.8**	**83.5**	**79.9**

CopyPasteLLM achieves 92.8% on FaithEval (+12.6 pp over Context-DPO) and 83.6%/80.9%/86.8% on ConFiQA vs Attributed's 51.5%/53.3%/37.3%, using only 365 training samples (≈1/50th baseline data).

4.3 Non-Counterfactual Evaluation Results

Table 2 evaluates performance when context aligns with parametric knowledge, testing whether Copy-Paste maintains effectiveness in standard RAG settings.

Table 2: Accuracy in non-counterfactual settings. PubMedQA is evaluated on artificial subset 20,000 samples (none used for CopyPasteLLM training, see Appendix Table **??**). ConFiQA uses Original context and Original answers.

Method	Mistral-7B-v0.2				Llama-3-8B				Llama-3.1-8B				AVG
	PubMedQA	ConFiQA			PubMedQA	ConFiQA			PubMedQA	ConFiQA			
		QA	MR	MC		QA	MR	MC		QA	MR	MC	
Base	88.60	96.22	71.20	72.27	97.3	98.02	93.00	91.02	**98.15**	97.93	89.48	89.97	90.26
CopyPasteLLM (Ours)	**91.40**	**97.43**	**91.87**	**91.20**	**97.5**	**99.30**	**97.17**	**96.27**	97.67	**99.02**	**94.95**	**94.92**	**95.73**

CopyPasteLLM improves average accuracy from 90.26% to 95.73% (+5.47 pp), with notable gains on PubMedQA (Mistral-7B: 91.4% vs 88.6%) and ConFiQA (+2.12/+11.82/+4.37 pp on QA/MR/MC).

4.4 Copy-Paste-Prompting Analysis

Table 3 compares our three Copy-Paste-Prompting methods against baseline approaches across different model sizes.

Stage 1 evaluation shows CP-Refine achieves best hallucination reduction (+10.9–19.1%), CP-Order excels in extractiveness, and CP-Link balances extractiveness and fluency. All Copy-Paste methods outperform baselines, with performance scaling by model size, validating our design choices.

4.5 Interpretability Analysis

To understand how CopyPasteLLM achieves improved performance, we conduct interpretability analysis using Context-Parameter Copying Capturing.

AI-Generated

Table 3: Performance comparison of Copy-Paste-Prompting against baselines across models and datasets. Methods with colored backgrounds are our proposed Copy-Paste-Prompting. **Bold** indicates the best performance, underlined indicates the second-best performance. *Faith.*: Faithfulness (*M.C.*: MiniCheck, *A.S.*: AlignScore), *Hallu.*: Hallucination, *Flu.*: Fluency.

Method	RAGTruth					FaithEval					PubmedQA					AVERAGE	
	Faith.		Hallu.		Flu.	Faith.		Hallu.		Flu.	Faith.		Hallu.		Flu.	Faith.	Hallu.
	M.C.	A.S.	Twist	Causal		M.C.	A.S.	Twist	Causal		M.C.	A.S.	Twist	Causal			
Mistral-7B-Instruct-v0.2 **(7B)**																	
Attributed	69.58	63.43	1506.9	1494.5	19.54	88.28	90.67	1527.1	1513.7	37.32	75.49	77.90	1464.7	1450.4	23.53	77.56	1492.9
Citations	57.82	49.39	1472.5	1475.7	**14.41**	73.50	74.25	1392.1	1416.2	27.98	55.79	52.35	1415.9	1370.0	**13.93**	60.52	1423.7
CP-Link	89.39	**75.45**	1518.9	1519.5	73.33	93.41	92.44	1510.9	1521.9	49.40	**96.50**	**88.52**	1518.4	**1580.7**	35.57	**89.29**	1528.4
CP-Order	**91.25**	71.98	1467.9	1472.4	65.62	**94.89**	92.27	1522.6	1501.5	43.74	93.18	82.35	1528.3	1559.1	32.65	87.65	1508.6
CP-Refine	82.18	74.56	**1533.8**	**1537.9**	18.46	92.85	**94.68**	**1547.4**	**1546.7**	**26.63**	91.52	88.21	**1572.7**	1539.7	17.79	87.33	**1546.4**
Llama-3.1-8B-Instruct **(8B)**																	
Attributed	57.02	65.29	1526.3	1554.3	26.22	85.22	85.65	1516.5	1536.9	330.8	71.10	60.01	1530.0	1553.1	47.36	70.72	1536.2
Citations	64.27	72.81	1428.5	**1574.4**	**16.78**	88.81	86.80	1486.2	**1555.6**	39.65	78.56	73.03	1403.4	1463.4	19.11	77.38	1485.3
CP-Link	70.58	78.83	1401.1	1328.3	17.83	91.54	89.23	1456.2	1366.3	**24.09**	80.74	80.79	1396.4	1371.1	19.65	81.95	1386.6
CP-Order	75.30	**94.81**	1498.4	1498.0	26.35	**95.44**	**98.12**	**1523.2**	1541.2	33.46	87.07	**97.62**	**1633.6**	**1559.1**	27.83	**91.39**	1542.3
CP-Refine	**77.30**	88.52	**1645.7**	1545.0	17.75	94.40	93.71	1517.9	1500.1	26.99	**87.29**	91.19	1536.5	1553.2	**18.64**	88.74	**1549.7**
Qwen2.5-72B-Instruct **(72B)**																	
Attributed	57.00	62.23	1504.5	1525.5	19.68	85.74	83.03	1537.3	1490.0	293.8	77.99	69.25	1509.9	1441.5	33.42	72.54	1501.5
Citations	74.32	77.52	1455.5	1498.0	**18.61**	90.98	88.30	1456.5	1476.7	34.67	82.01	76.62	1358.8	1413.6	22.89	81.63	1443.2
CP-Link	75.75	85.37	1446.3	1363.2	27.47	92.88	92.00	1443.5	1424.2	39.55	86.21	88.58	1527.9	1489.2	33.43	86.80	1449.1
CP-Order	76.32	**94.60**	1509.2	**1589.6**	30.56	**95.78**	**98.16**	**1539.3**	**1579.7**	38.11	87.85	**97.52**	1546.8	1575.9	35.26	**91.71**	**1556.8**
CP-Refine	**78.14**	90.88	**1584.6**	1523.7	20.12	94.72	95.48	1523.4	1529.4	**27.65**	**88.88**	95.04	**1556.7**	**1579.9**	**20.29**	90.52	1549.6
DeepSeek-V3-0324 **(671B)**																	
Attributed	56.42	59.60	1417.1	1449.1	27.52	86.90	83.46	1524.3	1535.0	63.27	75.56	69.24	1449.2	1487.9	36.88	71.86	1477.1
Citations	62.32	64.45	1510.8	1565.6	34.63	87.38	85.69	1463.0	1477.0	36.09	75.93	71.85	1460.4	1387.5	23.27	74.60	1477.4
CP-Link	70.59	72.54	1382.9	1360.3	34.19	92.60	88.08	1489.1	1374.8	35.55	81.56	77.67	1380.9	1351.1	28.54	80.51	1389.9
CP-Order	75.53	**92.87**	1579.4	1555.2	59.11	**95.23**	**97.79**	**1569.9**	1548.1	34.30	87.20	**97.38**	1561.8	1621.7	27.56	**91.00**	1572.7
CP-Refine	**77.14**	90.02	**1609.8**	**1569.7**	**22.57**	94.45	93.06	1453.7	**1565.2**	**33.84**	**87.39**	91.05	**1647.7**	**1651.7**	**21.91**	88.85	**1583.0**

Generated

4.5.1 Hidden State Analysis

Figure 3 shows dimensionality reduction visualization of hidden state distributions between contextual (CTX) and parametric (Para.) knowledge.

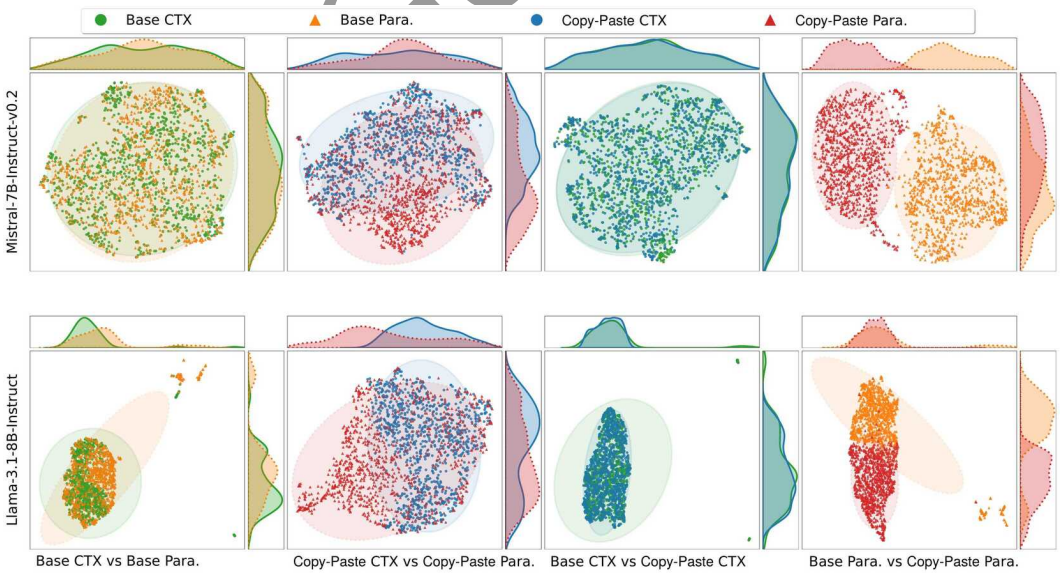

Figure 3: Hidden state distributions for contextual (CTX) and parametric (Para.) knowledge on PubMedQA. CopyPasteLLM shows clearer separation between CTX and Para. representations compared to base models, indicating stronger differentiation of knowledge sources.

CopyPasteLLM creates clearer separation between contextual and parametric knowledge representations, indicating stronger contextual prioritization.

4.5.2 LOGITS POWER ANALYSIS

Figure 4 examines logits power distribution across response lengths for contextual versus parametric knowledge.

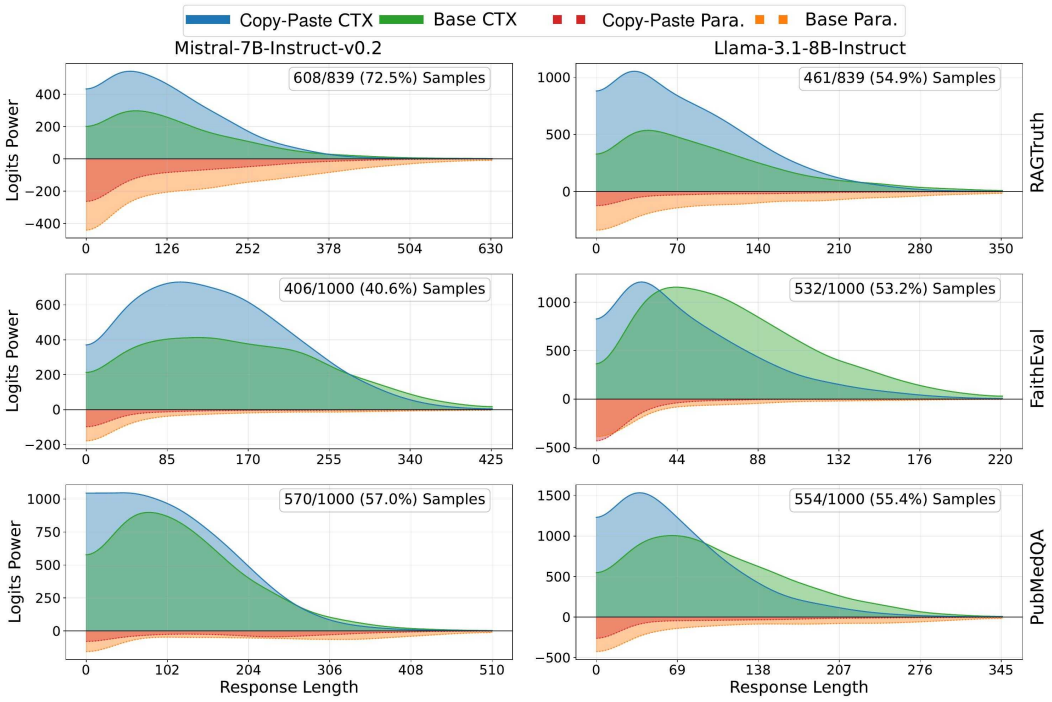

Figure 4: Logits power across response lengths for contextual (CTX) and parametric (Para.) knowledge. CopyPasteLLM shows increased CTX power and decreased Para. power compared to base models, indicating stronger reliance on contextual information.

CopyPasteLLM shows increased CTX logits power, suppressed parametric confidence, and consistent behavior across models, revealing how Copy-Paste training favors contextual over parametric knowledge.

4.6 ABLATION STUDIES

Ablation studies validate key design choices: optimal data size is 365 samples (diminishing returns beyond), extractiveness threshold of 0.99 balances performance and fluency, Elo-based preference construction improves accuracy by 15.3% over random pairing, and combining all three Copy-Paste methods maximizes preference diversity.

5 RELATED WORK

While Retrieval-Augmented Generation has emerged as a primary approach for grounding LLM responses (Gao et al., 2023), context faithfulness remains challenging as models prefer parametric over provided knowledge (Li et al., 2025). Various hallucination detection methods focus on post-hoc detection rather than prevention (Sun et al., 2024), while extractive generation ensures faithfulness through direct copying but sacrifices fluency (Zhang et al., 2023). Recent attribution approaches explore explicit citations (Gao et al., 2023), though they emphasize post-hoc rather than inherent grounding.

Our work differs by: (1) operationalizing faithfulness through explicit copying behavior, (2) providing data-efficient training with minimal samples, and (3) offering inherent attribution through direct text reuse for transparent verification.

6 CONCLUSION

We present Copy-Paste, a novel framework that mitigates LLM hallucinations through explicit copying behavior and contextual trust internalization. Our two-stage approach uses Copy-Paste-Prompting (CP-Order, CP-Link, CP-Refine) and CopyPasteLLM training via preference learning to internalize faithful generation patterns.

CopyPasteLLM achieves 92.8% accuracy using only 365 samples (≈1/50th baseline data), improving 12.2–24.5 pp over baselines. Non-counterfactual accuracy rises from 90.26% to 95.73%. Interpretability analysis shows suppressed parametric confidence and strengthened contextual reliance. Copy-Paste provides inherent attribution, data efficiency, and consistent improvements, making it valuable for high-stakes applications.

7 ETHICS STATEMENT

This work improves AI reliability by reducing hallucinations. Key considerations: users must evaluate source quality as copying cannot address biased context; direct copying may perpetuate source biases; the method suits factual over creative generation; inherent attribution enables verification. We encourage responsible deployment with appropriate oversight.

8 REPRODUCIBILITY STATEMENT

We provide complete reproducibility materials: source code, documentation, trained weights (Llama-3-8B, Mistral-7B), hyperparameters, CopyPasteSeed365 dataset, and evaluation scripts, enabling full reproduction and domain extension.

REFERENCES

Baolong Bi, Shaohan Huang, Yiwei Wang, Tianchi Yang, Zihan Zhang, Haizhen Huang, Lingrui Mei, Junfeng Fang, Zehao Li, Furu Wei, Weiwei Deng, Feng Sun, Qi Zhang, and Shenghua Liu. Context-DPO: Aligning language models for context-faithfulness. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (eds.), Findings of the Association for Computational Linguistics: ACL 2025, pp. 10280–10300, Vienna, Austria, July 2025. Association for Computational Linguistics. ISBN 979-8-89176-256-5. doi: 10.18653/v1/2025.findings-acl.536. URL https://aclanthology.org/2025.findings-acl.536/.

Tianyu Gao, Howard Yen, Jiatong Yu, and Danqi Chen. Enabling large language models to generate text with citations. In Houda Bouamor, Juan Pino, and Kalika Bali (eds.), Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, pp. 6465–6488, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.emnlp-main.398. URL https://aclanthology.org/2023.emnlp-main.398/.

Pengcheng Huang, Zhenghao Liu, Yukun Yan, Haiyan Zhao, Xiaoyuan Yi, Hao Chen, Zhiyuan Liu, Maosong Sun, Tong Xiao, Ge Yu, and Chenyan Xiong. Parammute: Suppressing knowledge-critical ffns for faithful retrieval-augmented generation. In Proceedings of the 39th International Conference on Neural Information Processing Systems, NIPS '25, Red Hook, NY, USA, 2025. Curran Associates Inc. URL https://neurips.cc/virtual/2025/poster/119254.

Yuepei Li, Kang Zhou, Qiao Qiao, Bach Nguyen, Qing Wang, and Qi Li. Investigating context faithfulness in large language models: The roles of memory strength and evidence style. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (eds.), Findings of ACL 2025, pp. 4789–4807, Vienna, Austria, July 2025. Association for Computational Linguistics. ISBN 979-8-89176-256-5.

Shuzheng Si, Haozhe Zhao, Cheng Gao, Yuzhuo Bai, Zhitong Wang, Bofei Gao, Kangyang Luo, Wenhao Li, Yufei Huang, Gang Chen, et al. Teaching large language models to maintain contextual faithfulness via synthetic tasks and reinforcement learning. arXiv preprint arXiv:2505.16483, 2025.

AI-generated

ZhongXiang Sun, Xiaoxue Zang, Kai Zheng, Jun Xu, Xiao Zhang, Weijie Yu, Yang Song, and Han Li. ReDeEP: Detecting hallucination in retrieval-augmented generation via mechanistic interpretability. In ICLR 2025 Spotlight, October 2024.

Santosh T.y.s.s, Youssef Tarek Elkhayat, Oana Ichim, Pranav Shetty, Dongsheng Wang, Zhiqiang Ma, Armineh Nourbakhsh, and Xiaomo Liu. CoCoLex: Confidence-guided copy-based decoding for grounded legal text generation. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (eds.), Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), pp. 19002–19018, Vienna, Austria, July 2025. Association for Computational Linguistics. ISBN 979-8-89176-251-0. URL https://aclanthology.org/2025.acl-long.931/.

Haopeng Zhang, Xiao Liu, and Jiawei Zhang. Extractive summarization via ChatGPT for faithful summary generation. In Houda Bouamor, Juan Pino, and Kalika Bali (eds.), Findings of the Association for Computational Linguistics: EMNLP 2023, pp. 3270–3278, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.findings-emnlp.214. URL https://aclanthology.org/2023.findings-emnlp.214/.

Wenxuan Zhou, Sheng Zhang, Hoifung Poon, and Muhao Chen. Context-faithful prompting for large language models. In Houda Bouamor, Juan Pino, and Kalika Bali (eds.), Findings of the Association for Computational Linguistics: EMNLP 2023, pp. 14544–14556, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.findings-emnlp.968. URL https://aclanthology.org/2023.findings-emnlp.968/.

AI-Generated