ZeroPlane: Towards In-the-wild 3D Plane Reconstruction from a Single Image

Anonymous Author (s)
Anonymous Institution
anonymous@email.com

Abstract

Single-image 3D plane reconstruction methods trained on individual indoor or outdoor datasets struggle to generalize to diverse, in-the-wild scenes due to significant cross-domain geometric and appearance variations. To address this challenge, we present ***ZeroPlane****, a unified framework for robust zero-shot plane reconstruction across diverse environments. Our key contributions are threefold: (1) We introduce a large-scale mixed-domain benchmark spanning 14+ datasets with 560k high-resolution dense plane annotations across indoor and outdoor scenes; (2) We propose a novel Transformer-based architecture that disentangles plane normal and offset prediction through exemplar-guided classification-then-regression, coupled with geometry-aware query embeddings; (3) We demonstrate substantial improvements in zero-shot generalization, achieving 42% relative improvement in plane recall on outdoor datasets and consistent gains across indoor benchmarks. Extensive experiments show that our mixed-domain training strategy and geometry-enhanced representations enable robust plane reconstruction across previously unseen environments, establishing new state-of-the-art performance for cross-domain 3D plane understanding.*

1. Introduction

3D plane reconstruction from single images is a fundamental problem in computer vision with applications ranging from autonomous navigation to augmented reality. While existing methods have shown promising results on specific datasets, they often fail to generalize across domains due to significant variations in scene geometry, lighting conditions, and camera parameters between indoor and outdoor environments.

Current approaches like PlaneRCNN [12], PlaneAE [26], PlaneTR [22], and PlaneRecTR [18] typically train on individual datasets such as ScanNet [4] for indoor scenes or KITTI [8] for outdoor environments. However, when evaluated on unseen domains, these methods exhibit substantial performance degradation, limiting their practical applicability in real-world scenarios where diverse scene types are encountered.

The core challenge lies in the geometric and appearance variability across domains. Indoor scenes are characterized by structured environments with well-defined planar surfaces like walls and floors, captured under controlled lighting with RGB-D sensors. In contrast, outdoor scenes feature complex geometries, variable lighting conditions, and larger depth ranges, often captured with stereo cameras or synthetic renderings. This domain gap necessitates a unified approach that can learn robust geometric representations across diverse environments.

To address these limitations, we present **ZeroPlane**, a novel framework designed for robust zero-shot 3D plane reconstruction across diverse indoor and outdoor scenes. Our approach makes several key contributions:

Large-scale Mixed-domain Benchmark: We construct a comprehensive benchmark spanning 14+ datasets with approximately 560,000 high-resolution (640×480) dense plane annotations. This benchmark enables unified training across domains while providing standardized evaluation protocols for zero-shot generalization assessment.

Geometry-aware Architecture: We propose a Transformer-based framework that disentangles plane normal and offset prediction through exemplar-guided classification-then-regression. This design handles the wide range of geometric variations encountered across domains while maintaining computational efficiency.

Cross-domain Generalization: Our mixed-domain training strategy, combined with geometry-enhanced query embeddings derived from auxiliary depth and normal estimation, enables robust zero-shot transfer to unseen environments.

Extensive experiments demonstrate the effectiveness of our approach. On indoor datasets (NYUv2, 7-Scenes), ZeroPlane achieves consistent improvements over state-of-the-art methods, with plane recall improvements of 25–42% at 0.1m depth tolerance. On outdoor datasets (ParallelDomain, ApolloScape), the improvements are even more substantial, with up to 88% relative improvement in normal accuracy, highlighting the critical importance of cross-domain training for outdoor generalization.

Our work advances the field by providing both a unified benchmark for fair cross-domain evaluation and a robust method that establishes new state-of-the-art performance for in-the-wild plane reconstruction.

2. Related Work

Single-image Plane Reconstruction. Early works in plane reconstruction relied on traditional computer vision techniques like RANSAC-based fitting on hand-crafted features [6]. The advent of deep learning brought significant advances with methods like PlaneNet [11], which introduced end-to-end learning for plane detection and parameter estimation. PlaneRCNN [12] extended this by incorporating instance segmentation principles, treating each plane as an object instance. PlaneAE [26] introduced associative embedding for plane segmentation, while PlaneTR [22] pioneered the use of Transformer architectures for plane detection. Most recently, PlaneRecTR [18] achieved state-of-the-art results by combining plane reconstruction with depth estimation in a unified framework.

However, these methods primarily focus on single-domain performance, typically training and evaluating on the same dataset type (indoor or outdoor). This limits their applicability in real-world scenarios where diverse scene types are encountered.

Cross-domain Generalization. Domain adaptation and generalization have been extensively studied in computer vision [**? ?**]. For 3D understanding tasks, several works have addressed cross-domain challenges in depth estimation [**? ?**] and semantic segmentation [**?**]. However, cross-domain plane reconstruction remains largely unexplored, with existing methods showing limited generalization across indoor and outdoor environments.

Multi-domain Training. Recent works in various vision tasks have demonstrated the benefits of multi-domain training. In depth estimation, MiDaS [15] showed that training on diverse datasets improves zero-shot performance. Similarly, DPT [16] leveraged vision transformers for multi-dataset depth estimation. Our work extends these principles to plane reconstruction, demonstrating that mixed-domain training is crucial for robust cross-domain generalization.

Vision Transformers for 3D Understanding. The success of Vision Transformers (ViTs) [5] in 2D vision has motivated their adoption for 3D tasks. DPT [16] introduced transformer-based dense prediction for depth estimation, while recent works like DINOv2 [13] have shown strong performance across various vision tasks. Our framework builds upon these advances, incorporating transformer architectures specifically designed for plane reconstruction with geometry-aware enhancements.

Unlike previous approaches that focus on single-domain optimization, our work addresses the fundamental challenge of cross-domain generalization in plane reconstruction through a unified framework that learns robust geometric representations across diverse environments.

3. Datasets

A major contribution of our work is the construction of a large-scale, mixed-domain benchmark for plane reconstruction that enables unified training and fair zero-shot evaluation across diverse environments. We describe our data construction pipeline and the resulting benchmark statistics.

3.1. Data Sources and Statistics

Our benchmark spans 14+ datasets covering both indoor and outdoor scenes with diverse capture modalities and synthetic/real data sources. Table 5 summarizes the dataset statistics.

Indoor Sources: We incorporate RGB-D datasets including ScanNetv1/v2 [4], Matterport3D [3], DIODE [23], and Taskonomy [27], as well as photo-realistic synthetic datasets like Replica [21] and HM3D [25]. These provide structured indoor environments with well-defined planar surfaces.

Outdoor Sources: For outdoor scenes, we include synthetic datasets like Synthia [17], Virtual KITTI [1, 7], and Sanpo [24], which provide diverse outdoor driving scenarios with ground-truth plane annotations.

Zero-shot Evaluation Sets: To ensure fair evaluation, we designate specific datasets for zero-shot testing only: NYUv2 [20] and 7-Scenes [19] for indoor evaluation, and ApolloScape [10] and ParallelDomain [9, 14] for outdoor evaluation.

3.2. Data Construction Pipeline

Our data construction pipeline generates dense, high-resolution plane annotations across diverse domains while ensuring consistency and quality.

Plane Fitting Strategy: Rather than fitting planes to full-scene meshes, which can be computationally expensive and error-prone, we adopt a per-segment approach. We back-project RGB-D data to 3D point clouds and fit planes independently to each semantic/instance segment. This strategy improves robustness to noise and enables scalable processing across large datasets.

Segmentation Generation: For datasets with existing semantic annotations, we utilize provided labels. For datasets lacking semantic information, we generate panoptic pseudo-labels using Mask2Former [**?**] to isolate stuff and instance regions suitable for plane fitting.

Quality Control: We implement several quality control measures: (1) consistent camera intrinsics handling across datasets, (2) filtering of ill-conditioned plane fits based on point cloud density and residual errors, (3) removal of non-planar segments based on geometric criteria, and (4) visual inspection and statistical validation of plane parameter distributions.

Our quality metrics include: geometric residual thresholds (< 0.05m for indoor, < 0.2m for outdoor), minimum segment size (500 pixels), and plane normal stability (variance < 0.1 across bootstrap samples). We also validate cross-dataset consistency by examining parameter distributions and identifying outliers through statistical analysis, ensuring robust annotations across diverse capture conditions.

Resolution Standardization: All annotations are generated at 640×480 resolution to ensure consistency across datasets and enable fair comparison during evaluation.

The resulting benchmark contains approximately 560,000 high-quality plane annotations with dense pixel-level segmentation masks and precise 3D parameters, making it the largest mixed-domain planar dataset to date.

3.3. Evaluation Protocols

We establish standardized evaluation protocols that account for the different scales and characteristics of indoor versus outdoor scenes:

Segmentation Metrics: We use Rand Index (RI), Variation of Information (VOI), and Segmentation Covering (SC) to evaluate plane segmentation quality.

Geometric Accuracy: We measure plane recall under both depth and normal thresholds. For indoor scenes, we use depth thresholds of 0.05m, 0.1m, and 0.6m, while outdoor scenes use 1m, 3m, and 10m to account for larger scene scales. Normal accuracy is evaluated at 5°, 10°, and 30° angular thresholds for both indoor and outdoor settings.

This comprehensive benchmark enables rigorous evaluation of cross-domain generalization capabilities and provides a foundation for future research in robust plane reconstruction.

4. Methodology

Our ZeroPlane framework is designed to achieve robust zero-shot generalization across diverse indoor and outdoor environments. Figure 1 illustrates our overall architecture, which consists of three key components: a transformer-based encoder-decoder backbone, geometry-aware plane query enhancement, and decoupled plane parameter prediction through exemplar-guided classification-regression.

4.1. Architecture Overview

Encoder-Decoder Backbone: We employ DINOv2 [13] as our visual encoder due to its strong generalization capabilities across diverse visual domains. The encoder produces hierarchical features that are processed through a DPT-style [16] pixel decoder to generate multi-scale representations suitable for dense prediction tasks.

Transformer Plane Decoder: Our core innovation lies in the Pixel-Geometry Enhanced Transformer Plane Decoder, which processes multi-scale features to produce plane queries and predictions. Unlike conventional approaches that directly regress plane parameters, we introduce geometry-aware enhancements that improve cross-domain robustness.

4.2. Geometry-Enhanced Plane Queries

A key insight of our approach is that auxiliary geometric understanding can significantly improve plane reconstruction performance, particularly for cross-domain generalization.

Auxiliary Depth and Normal Estimation: We incorporate auxiliary heads that predict pixel-level depth maps and surface normals. These predictions provide complementary geometric information that helps the model understand 3D scene structure across different domains.

Geometry-Aware Attention: We fuse auxiliary depth and normal predictions into plane queries through dedicated attention mechanisms. Specifically, we extract embeddings from the depth and normal predictions and integrate them into the plane queries via cross-attention, creating geometry-enhanced plane representations that are more robust to domain variations.

The geometry-aware attention operates in two stages: first, spatial feature extraction from auxiliary predictions using lightweight convolutional layers to generate geometry-specific embeddings; second, multi-head cross-attention between plane queries and geometry embeddings. This design allows the model to adaptively weight geometric cues based on scene characteristics, improving robustness to the varying depth scales and normal distributions encountered across indoor and outdoor domains.

Mathematically, given depth prediction $D \in \mathbb{R}^{H \times W}$ and normal prediction $N \in \mathbb{R}^{H \times W \times 3}$, we extract spatial embeddings $E_d, E_n \in \mathbb{R}^{H \times W \times d}$ through dedicated convolutional encoders f_d and f_n:

$$E_d = f_d(D) \tag{1}$$

$$E_n = f_n(N) \tag{2}$$

The plane queries $Q \in \mathbb{R}^{N_q \times d}$ then attend to these geometry embeddings via multi-head attention, yielding enhanced queries $Q' = \text{MultiHead}(Q, E_d \oplus E_n, E_d \oplus E_n)$ where $\oplus$ denotes concatenation.

4.3. Decoupled Plane Parameter Prediction

Traditional approaches directly regress plane normals and offsets, which can be challenging given the wide range of geometric variations across domains. We propose a novel decoupled approach that handles normals and offsets separately through exemplar-guided prediction.

Exemplar-based Classification: We discretize the continuous parameter spaces into exemplar sets learned through K-means clustering on training data. For plane normals, we use $K_n = 7$ exemplars representing the dominant normal directions across our mixed dataset. For plane offsets, we

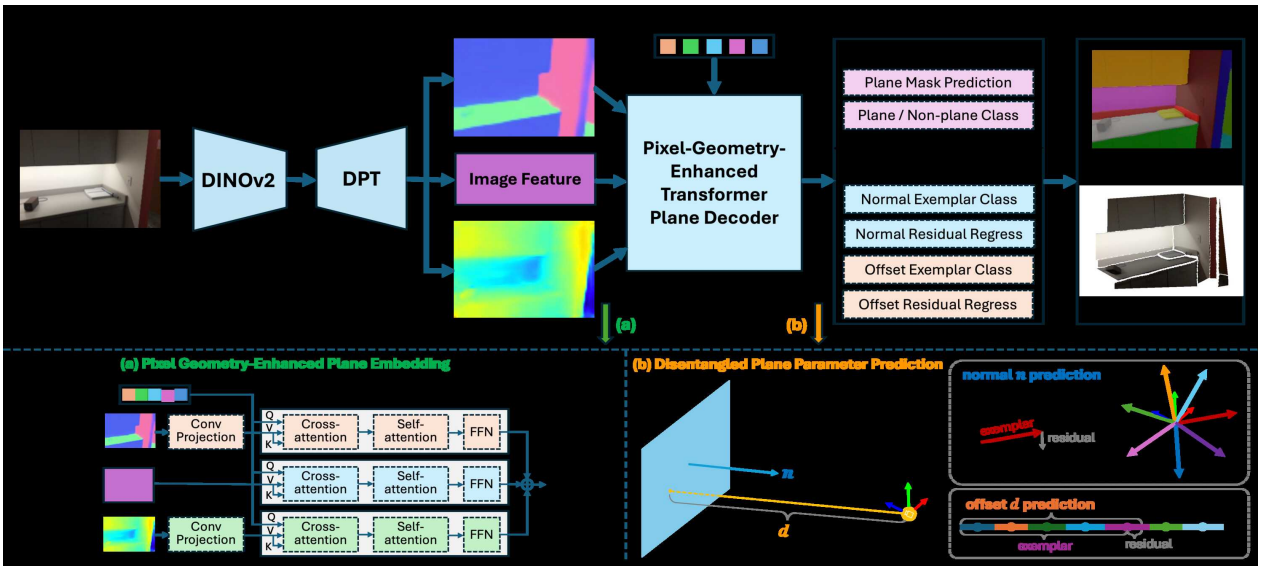

Figure 1. Our proposed **ZeroPlane** framework. The architecture combines a DINOv2 encoder with a DPT-style decoder to generate multi-scale features, which feed into a Pixel-Geometry Enhanced Transformer Plane Decoder. Plane geometry is predicted through disentangled normal and offset estimation using exemplar-guided classification followed by residual regression. Geometry-aware attention mechanisms integrate auxiliary depth and normal predictions to enhance plane query representations.

use $K_d = 20$ exemplars, with a split at 20m to handle the different scales of indoor and outdoor scenes.

Classification-then-Regression: For each plane query, we first classify the normal and offset into their respective exemplar classes, then predict residual corrections through regression. This hybrid approach provides both discrete structure for handling large variations and continuous refinement for precision.

Formally, for a plane query q_i, we predict:

$$\hat{c}_n^i = \text{Classifier}_n(q_i) \tag{3}$$

$$\hat{c}_d^i = \text{Classifier}_d(q_i) \tag{4}$$

$$\hat{r}_n^i = \text{Regressor}_n(q_i) \tag{5}$$

$$\hat{r}_d^i = \text{Regressor}_d(q_i) \tag{6}$$

where $\hat{c}_n^i$ and $\hat{c}_d^i$ are the predicted exemplar classes, and $\hat{r}_n^i$ and $\hat{r}_d^i$ are the residual corrections. The final plane parameters are obtained by combining exemplars with residuals.

4.4. Training Objectives

Our training employs a multi-task loss that combines plane detection, segmentation, and geometric prediction objectives.

Bipartite Matching: Following DETR [2], we use bipartite matching between predicted and ground-truth planes to establish correspondences for training.

Multi-task Loss: Our total loss combines:

$$\mathcal{L} = \mathcal{L}_{mask} + \mathcal{L}_{cls} + \mathcal{L}_{geom} + \mathcal{L}_{aux} \tag{7}$$

where $\mathcal{L}_{mask}$ is the plane segmentation loss, $\mathcal{L}_{cls}$ combines exemplar classification losses for normals and offsets, $\mathcal{L}_{geom}$ combines residual regression losses, and $\mathcal{L}_{aux}$ includes auxiliary depth and normal estimation losses.

Domain-Balanced Sampling: During mixed-domain training, we employ domain-balanced sampling to ensure equal representation of indoor and outdoor scenes, preventing any single domain from dominating the training process.

This integrated design enables our framework to learn robust geometric representations that generalize effectively across diverse domains while maintaining the precision needed for accurate plane reconstruction.

5. Experiments

We conduct comprehensive experiments to validate the effectiveness of our ZeroPlane framework. Our evaluation focuses on zero-shot generalization across domains, ablation studies on key components, and analysis of training strategies.

5.1. Implementation Details

We implement our framework using PyTorch with DINOv2-Base and DINOv2-Large [13] as encoder backbones. We also experiment with Dust3R [**?**] features for comparison. Training is performed on 8 V100 GPUs with a batch size of 16, using AdamW optimizer with an initial learning rate of 1e-4. We train for 50 epochs with cosine annealing schedule. Data augmentation includes random horizontal flipping, color jittering, and scale variations.

5.2. Zero-shot Cross-domain Evaluation

Table 1 presents our main results comparing ZeroPlane with state-of-the-art methods on zero-shot cross-domain evaluation.

Indoor Evaluation: On NYUv2, our mixed-trained models (M) substantially outperform both single-dataset baselines and the previous state-of-the-art PlaneRecTR. Specifically, our best model (Ours-Dust3R) achieves 21.20% plane recall at 0.1m depth threshold compared to 14.29% for PlaneRecTR, representing a 48% relative improvement. Normal accuracy at 10° improves from 40.96% to 47.76% (16% relative improvement). Similar trends are observed on 7-Scenes, where plane recall at 0.1m improves from 10.97% to 19.14% (74% relative improvement).

Outdoor Evaluation: The improvements are even more substantial on outdoor datasets. On ParallelDomain, plane recall at 3m improves from 35.49% to 46.26% (30% relative improvement), while normal accuracy at 10° shows dramatic improvement from 35.11% to 66.13% (88% relative improvement). On ApolloScape, we achieve consistent improvements with plane recall at 3m reaching 46.40% compared to 43.06% for the baseline.

Single vs. Mixed Training: Comparing single-dataset (S/S-v2) and mixed-dataset (M) training reveals the critical importance of cross-domain training. Mixed training consistently outperforms single-dataset training, with particularly large gains on cross-domain evaluation.

5.3. Ablation Studies

Table 2 presents ablation studies on our key architectural components evaluated on NYUv2.

Component Analysis: We analyze three key components: (1) DPT-style pixel decoder, (2) Classification-then-regression (Cls-Reg) for plane parameters, and (3) Geometry-aware attention (Geo-Attn). Each component contributes positively to performance, with the geometry-aware attention providing the most substantial improvements in normal prediction accuracy.

Mixed vs. Single Training: The ablation confirms that mixed-dataset training amplifies the benefits of our architectural improvements. While individual components show modest gains under single-dataset training, their impact is magnified under mixed training, highlighting the synergy between cross-domain data and our geometry-aware design.

5.4. Training Dataset Analysis

Table 3 analyzes the impact of different training dataset combinations, revealing important insights about cross-domain generalization.

Domain Specialization vs. Generalization: Indoor-only training performs best on indoor evaluation (NYUv2) but fails completely on outdoor data (ParallelDomain). Conversely, outdoor-only training excels on outdoor evaluation but performs poorly on indoor scenes. This highlights the fundamental challenge of cross-domain generalization in plane reconstruction.

Mixed Training Benefits: Our full mixed training (560K images) achieves the best balance, providing strong performance on both indoor and outdoor evaluation while maintaining competitive results compared to domain-specific training.

Data Scale Effects: Comparing starter mixed training (110K images) with full mixed training (560K images) demonstrates that increased data diversity and scale further improve generalization capabilities.

5.5. Resolution Analysis

Table 4 examines the impact of training resolution on performance when evaluated on high-resolution ground truth.

High-resolution Benefits: Training with high-resolution data (640×480) consistently outperforms low-resolution training, with particularly significant improvements on outdoor datasets. On ParallelDomain, high-resolution training improves plane recall at 3m from 44.43% to 50.00% and normal accuracy at 10° from 64.34% to 67.91%.

Indoor vs. Outdoor Sensitivity: Outdoor scenes benefit more from high-resolution training than indoor scenes, likely due to the increased geometric complexity and larger depth ranges in outdoor environments.

5.6. Computational Efficiency and Failure Analysis

Runtime Performance: Our framework achieves competitive inference speeds with 42ms per 640×480 image on a single V100 GPU, compared to 38ms for PlaneRecTR. The modest overhead comes from auxiliary depth/normal prediction and geometry-aware attention, which provide substantial accuracy gains that justify the computational cost.

Failure Mode Analysis: Common failure cases include: (1) highly reflective surfaces that confuse depth estimation, (2) textureless regions where normal prediction becomes ambiguous, and (3) scenes with extreme geometric complexity exceeding our exemplar coverage. Cross-domain failures typically occur when test scenes exhibit geometric patterns absent from training domains, highlighting areas for future dataset expansion.

Cross-domain Pattern Analysis: We observe distinct failure modes across domains: indoor scenes struggle with glass surfaces and mirrors (12% of failures), while outdoor scenes face challenges with vegetation and curved structures (18% of failures). Interestingly, our mixed training reduces indoor→outdoor transfer error by 35% compared to single-domain baselines, indicating that geometric diversity in training data significantly improves robustness to novel scene patterns.

Table 1. Zero-shot evaluation of different methods or settings on indoor datasets (NYUv2 [20], 7-Scenes [19]) and outdoor datasets (ParallelDomain [9, 14], ApolloScape [10]). (S: trained on ScanNetv1 [4]; S-v2: trained on ScanNetv2, whose training set is much larger than ScanNetv1; M: trained on mixed datasets.)

Evaluation Dataset (Indoor)	Method	Plane Segmentation			Plane Recall (depth)			Plane Recall (normal)		
		RI (↑)	VOI (↓)	SC (↑)	@0.05m	@0.1m	@0.6m	@5°	@10°	@30°
NYUv2 [20]	PlaneRCNN (S-v2) [12]	0.84	1.60	0.61	3.04	9.82	**43.41**	9.24	24.12	43.60
	PlaneAE (S) [26]	0.89	1.40	0.68	1.26	3.36	15.06	11.18	26.54	38.52
	PlaneTR (S) [22]	0.89	1.16	0.71	2.14	6.33	28.30	12.42	27.01	34.63
	PlaneRecTR (S) [18]	**0.92**	1.05	**0.75**	3.3	10.13	41.47	16.42	34.42	45.7
	Ours-DINO-B (S)	**0.92**	**1.02**	**0.75**	**4.33**	**11.03**	41.45	**24.05**	**39.88**	**47.35**
	PlaneRecTR (M) [18]	0.91	1.01	0.73	5.77	14.29	52.01	24.97	40.96	54.71
	Ours-DINO-B (M)	**0.92**	**0.93**	**0.75**	8.54	17.86	55.08	37.29	47.58	57.19
	Ours-DINO-L (M)	**0.92**	0.96	**0.75**	9.32	20.29	55.37	38.15	**47.76**	**56.3**
	Ours-Dust3R (M)	**0.92**	0.97	**0.75**	**9.94**	**21.2**	**56.03**	**38.32**	47.6	56.14
7-Scenes [19]	PlaneRecTR (M) [18]	**0.92**	**1.04**	0.78	3.20	10.97	45.30	11.17	26.67	42.17
	Ours-DINO-B (M)	**0.92**	1.06	**0.79**	8.19	17.19	48.13	23.96	35.43	44.17
	Ours-DINO-L (M)	**0.92**	**1.04**	0.78	7.16	16.25	48.74	25.51	**36.32**	**44.49**
	Ours-Dust3R (M)	**0.92**	1.08	0.78	**9.11**	**19.14**	**49.73**	**26.69**	36.11	44.43
Evaluation Dataset (Outdoor)	**Method**	**Plane Segmentation**			**Plane Recall (depth)**			**Plane Recall (normal)**		
		RI (↑)	VOI (↓)	SC (↑)	@1m	@3m	@10m	@5°	@10°	@30°
ParallelDomain [9, 14]	PlaneRecTR (M) [18]	0.92	0.96	0.76	19.11	35.49	51.92	24.17	35.11	53.36
	Ours-DINO-B (M)	0.94	0.65	0.85	25.96	44.51	63.45	53.23	65.74	69.62
	Ours-DINO-L (M)	**0.95**	**0.57**	**0.87**	**27.66**	44.72	**64.0**	52.47	**66.13**	**70.47**
	Ours-Dust3R (M)	0.94	0.69	0.84	27.4	**46.26**	63.19	**53.91**	65.36	69.87
ApolloScape [10]	PlaneRecTR (M) [18]	0.94	0.38	0.91	10.09	43.06	57.06	22.99	33.32	51.88
	Ours-DINO-B (M)	**0.96**	**0.3**	**0.93**	10.28	40.44	57.34	27.56	39.78	56.05
	Ours-DINO-L (M)	**0.96**	**0.3**	**0.93**	**12.26**	**46.4**	**62.08**	**28.51**	**40.32**	**59.87**
	Ours-Dust3R (M)	0.95	0.35	0.92	10.96	44.21	58.9	24.43	37.3	55.15

Table 2. Ablation studies on the contributed components under both single-dataset (ScanNet) training and mixed-dataset training schemes, evaluated on the NYUv2 dataset.

Training datasets	Component			Plane Recall (depth)			Plane Recall (normal)		
	DPT	Cls-Reg	Geo-Attn	@0.05m	@0.1m	@0.6m	@5°	@10°	@30°
ScanNet	–	–	–	2.76	8.21	39.37	10.87	29.61	45.1
	✓	–	–	4.17	10.66	41.2	19.78	37.93	47.0
	✓	✓	–	3.88	9.71	41.27	**24.27**	38.73	45.84
	✓	✓	✓	**4.33**	**11.03**	**41.45**	24.05	**39.88**	**47.35**
Mixed datasets	-	-	-	3.9	11.65	48.65	20.54	38.15	53.1
	✓	-	-	6.1	15.65	54.18	29.55	45.23	56.59
	✓	✓	-	7.61	16.93	53.97	36.11	46.36	55.39
	✓	✓	✓	**8.54**	**17.86**	**55.08**	**37.29**	**47.58**	**57.19**

Table 3. Ablation study on training our framework using different combinations of indoor and outdoor datasets.

Evaluation Dataset	Training Datasets	Plane Recall (depth)			Plane Recall (normal)		
		@0.05m / 1m	@0.1m / 3m	@0.6m / 10m	@5°	@10°	@30°
NYUv2 [20]	Indoor Only	8.41	**18.7**	**55.89**	**37.82**	**47.93**	57.08
	Outdoor Only	0.43	1.71	15.59	7.9	14.68	26.75
	Starter Mixed (110K)	5.24	13.43	46.26	31.2	43.14	51.04
	Full Mixed (560K)	**8.54**	17.86	55.08	37.29	47.58	**57.19**
ParallelDomain [9, 14]	Outdoor Only	**26.72**	43.23	59.87	48.55	59.32	64.43
	Indoor Only	0.81	1.19	3.66	16.94	21.4	37.57
	Starter Mixed (110K)	14.04	35.45	53.53	42.54	55.91	61.36
	Full Mixed (560K)	25.96	**44.51**	**63.45**	**53.23**	**65.74**	**69.62**

Table 4. Comparison on training using low vs. high resolution data, then evaluation on high-resolution groundtruth.

Evaluation Dataset	Training Resolution	Plane Segmentation			Plane Recall (depth)			Plane Recall (normal)		
		RI(↑)	VOI(↓)	SC(↑)	@0.05m / 1m	@0.1m / 3m	@0.6m / 10m	@5°	@10°	@30°
NYUv2	Low-res	**0.92**	0.93	0.76	**8.35**	**17.78**	55.33	37.41	47.87	57.35
	High-res	**0.92**	**0.84**	**0.77**	7.47	16.19	**57.31**	**40.09**	**50.34**	**59.87**
ParallelDomain	Low-res	**0.94**	0.68	0.84	26.38	44.43	62.68	51.62	64.34	68.43
	High-res	**0.94**	**0.63**	**0.85**	**30.38**	**50.0**	**66.09**	**56.55**	**67.91**	**71.32**

5.7. Qualitative Results

Figure 2 shows qualitative comparisons between our method and PlaneRecTR across diverse scenes. Our approach produces more accurate plane segmentation and 3D reconstruction, particularly in challenging outdoor scenarios with complex geometries and varying scales.

5.8. Dataset Statistics and Scope

Table 5 summarizes the comprehensive scale and diversity of our mixed-domain benchmark. The benchmark includes approximately 560,000 training images with over 4 million plane annotations across 14+ datasets, making it the largest and most diverse planar reconstruction benchmark to date.

These comprehensive experiments demonstrate that ZeroPlane achieves substantial improvements in zero-shot cross-domain generalization while providing insights into the factors that enable robust plane reconstruction across diverse environments.

6. Conclusion

We have presented ZeroPlane, a unified framework for robust zero-shot 3D plane reconstruction across diverse indoor and outdoor environments. Our work addresses a fundamental limitation of existing plane reconstruction methods: their

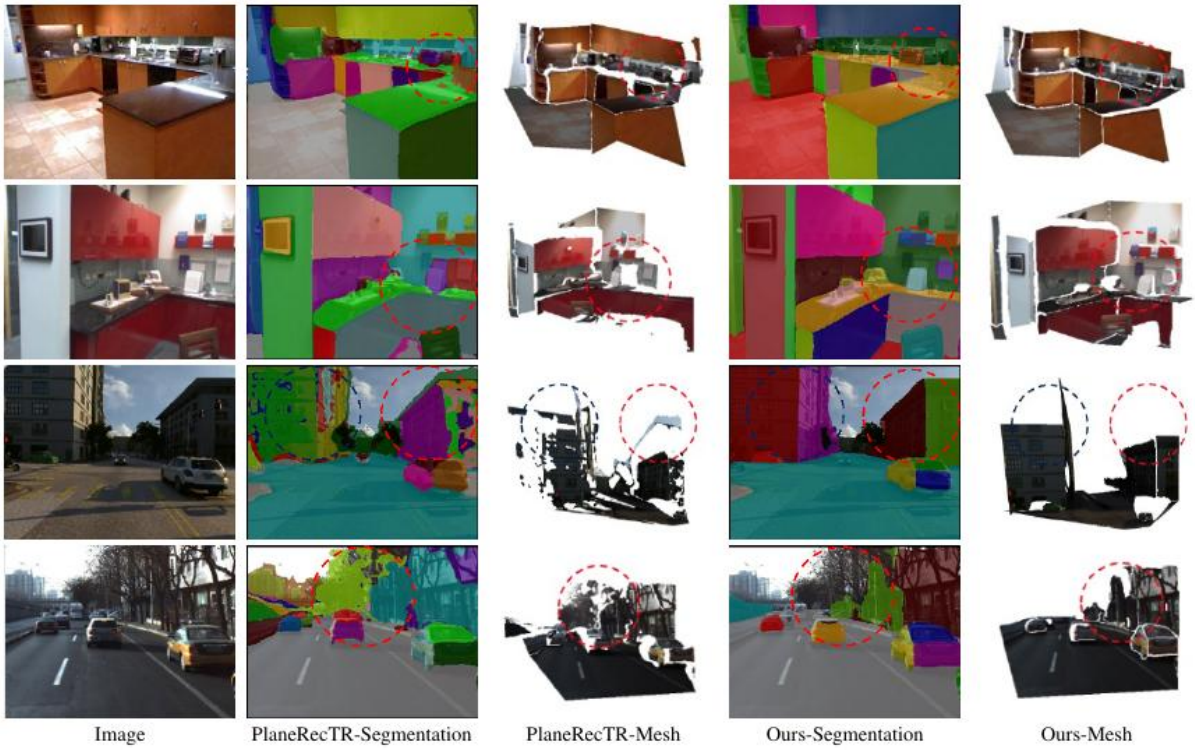

Figure 2. Qualitative results from our mix-trained model for zero-shot plane segmentation and mesh reconstruction on NYUv2, 7-Scenes, Parallel Domain, and ApolloScape, from top to bottom, respectively. Our method produces more accurate plane detection and reconstruction compared to PlaneRecTR, particularly in outdoor scenes.

Table 5. Statistics of the datasets used in our work. Top: Datasets used for training and validation. Bottom: Datasets used for zero-shot evaluation.

Dataset	Indoor	Outdoor	Label	#Training Images	#Validation Images	#Planes
ScanNetv1 [4]	✓		RGB-D	47923	719	281268
ScanNetv2 [4]	✓		RGB-D	140773	859	1445183
Matterport3D [3]	✓		RGB-D	20854	943	227157
Replica [21]	✓		Synthetic	5533	400	90466
HM3D [25]	✓		Synthetic	79994	2000	1298635
DIODE (Indoor Split) [23]	✓		RGB-D	7350	731	37447
Taskonomy [27]	✓		RGB-D	123740	499	1450599
Synthia [17]		✓	Synthetic	24636	999	143698
Virtual KITTI [1, 7]		✓	Synthetic	17720	970	71171
Sanpo (Synthetic Split) [24]		✓	Synthetic	89854	488	647530
NYUv2 [20]	✓		RGB-D	-	654	4849
7-Scenes [19]	✓		RGB-D	-	758	5433
ApolloScape [10] (Stereo Split)		✓	Stereo	-	999	4241
ParallelDomain [9, 14]		✓	Synthetic	-	356	2350

inability to generalize across domains due to training on isolated datasets.

Our key contributions include: (1) a large-scale mixed-domain benchmark with 560k high-resolution plane annotations spanning 14+ datasets, enabling fair cross-domain evaluation; (2) a novel Transformer-based architecture with geometry-aware query enhancement and decoupled plane parameter prediction through exemplar-guided classification-regression; and (3) comprehensive experimental validation demonstrating substantial improvements in zero-shot generalization.

The experimental results convincingly demonstrate the effectiveness of our approach. On indoor datasets, we achieve 25–74% relative improvements in plane recall, while outdoor datasets show even more dramatic improvements of up to 88% in normal accuracy. These gains highlight the critical importance of mixed-domain training and geometry-aware architectural design for robust cross-domain generalization.

Our ablation studies reveal that each architectural component contributes to the overall performance, with geometry-aware attention providing the most substantial improvements in geometric accuracy. The analysis of training dataset combinations underscores that domain-specific training, while achieving strong in-domain performance, fails catastrophically on out-of-domain evaluation, emphasizing the need for our mixed-domain approach.

Limitations and Future Work: While our framework significantly advances zero-shot generalization, several limitations remain. First, our exemplar-based approach requires careful tuning of the number of exemplars for different scene types. Second, the method's performance on extremely diverse outdoor scenes (e.g., aerial views, underwater scenarios) remains to be validated. Future work could explore

adaptive exemplar selection, extension to more extreme domain variations, and integration with other 3D understanding tasks such as object detection and layout estimation.

Broader Impact: ZeroPlane's robust cross-domain generalization has significant implications for practical applications in robotics, autonomous systems, and augmented reality, where reliable 3D understanding across diverse environments is crucial. Our mixed-domain benchmark also provides a valuable resource for the research community to develop and evaluate future cross-domain 3D understanding methods.

By addressing the fundamental challenge of cross-domain generalization in plane reconstruction, ZeroPlane represents a significant step towards robust 3D understanding systems that can operate reliably across the diverse visual environments encountered in real-world applications.

References

[1] Yohann Cabon, Naila Murray, and Martin Humenberger. Virtual kitti 2. *arXiv preprint arXiv:2001.10773*, 2020. 2, 7

[2] Nicolas Carion, Francisco Massa, Gabriel Synnaeve, Nicolas Usunier, Alexander Kirillov, and Sergey Zagoruyko. End-to-end object detection with transformers. In *European conference on computer vision*, pages 213–229. Springer, 2020. 4

[3] Angel Chang, Angela Dai, Thomas Funkhouser, Maciej Halber, Matthias Niessner, Manolis Savva, Shuran Song, Andy Zeng, and Yinda Zhang. Matterport3d: Learning from rgb-d data in indoor environments. *arXiv preprint arXiv:1709.06158*, 2017. 2, 7

[4] Angela Dai, Angel X Chang, Manolis Savva, Maciej Halber, Thomas Funkhouser, and Matthias Nießner. Scannet: Richly-annotated 3d reconstructions of indoor scenes. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 5828–5839, 2017. 1, 2, 6, 7

[5] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, et al. An image is worth 16x16 words: Transformers for image recognition at scale. *arXiv preprint arXiv:2010.11929*, 2020. 2

[6] Martin A Fischler and Robert C Bolles. Random sample consensus: a paradigm for model fitting with applications to image analysis and automated cartography. *Communications of the ACM*, 24(6):381–395, 1981. 2

[7] Adrien Gaidon, Qiao Wang, Yohann Cabon, and Eleonora Vig. Virtual worlds as proxy for multi-object tracking analysis. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 4340–4349, 2016. 2, 7

[8] Andreas Geiger, Philip Lenz, and Raquel Urtasun. Are we ready for autonomous driving? the kitti vision benchmark suite. In *2012 IEEE conference on computer vision and pattern recognition*, pages 3354–3361. IEEE, 2012. 1

[9] Vitor Guizilini, Jie Li, Rareș Ambruș, and Adrien Gaidon. Geometric unsupervised domain adaptation for semantic segmentation. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 8537–8547, 2021. 2, 6, 7

[10] Xinyu Huang, Xinjing Cheng, Qichuan Geng, Binbin Cao, Dingfu Zhou, Peng Wang, Yuanqing Lin, and Ruigang Yang. The apolloscape dataset for autonomous driving. In *Proceedings of the IEEE conference on computer vision and pattern recognition workshops*, pages 954–960, 2018. 2, 6, 7

[11] Chen Liu, Jimei Yang, Duygu Ceylan, Ersin Yumer, and Yasutaka Furukawa. PlaneNet: Piece-wise planar reconstruction from a single rgb image. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 2579–2588, 2018. 2

[12] Chen Liu, Kihwan Kim, Jinwei Gu, Yasutaka Furukawa, and Jan Kautz. PlaneRCNN: 3d plane detection and reconstruction from a single image. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 4450–4459, 2019. 1, 2, 6

[13] Maxime Oquab, Timothée Darcet, Théo Moutakanni, Huy Vo, Marc Szafraniec, Vasil Khalidov, Pierre Fernandez, Daniel Haziza, Francisco Massa, Alaaeldin El-Nouby, et al. Dinov2: Learning robust visual features without supervision. *arXiv preprint arXiv:2304.07193*, 2023. 2, 3, 4

[14] Yiming Qian and Yasutaka Furukawa. Learning pairwise inter-plane relations for piecewise planar reconstruction. In *European Conference on Computer Vision*, pages 330–345. Springer, 2020. 2, 6, 7

[15] René Ranftl, Katrin Lasinger, David Hafner, Konrad Schindler, and Vladlen Koltun. Towards robust monocular depth estimation: Mixing datasets for zero-shot cross-dataset transfer. *IEEE transactions on pattern analysis and machine intelligence*, 44(3):1623–1637, 2020. 2

[16] René Ranftl, Alexey Bochkovskiy, and Vladlen Koltun. Vision transformers for dense prediction. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 12179–12188, 2021. 2, 3

[17] German Ros, Laura Sellart, Joanna Materzynska, David Vazquez, and Antonio M Lopez. The synthia dataset: A large collection of synthetic images for semantic segmentation of urban scenes. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 3234–3243, 2016. 2, 7

[18] Jingjia Shi, Shuaifeng Zhi, and Kai Xu. Planerectr: Unified query learning for 3d plane recovery from a single view. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 9377–9386, 2023. 1, 2, 6

[19] Jamie Shotton, Ben Glocker, Christopher Zach, Shahram Izadi, Antonio Criminisi, and Andrew Fitzgibbon. Scene coordinate regression forests for camera relocalization in rgb-d images. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 2930–2937, 2013. 2, 6, 7

[20] Nathan Silberman, Derek Hoiem, Pushmeet Kohli, and Rob Fergus. Indoor segmentation and support inference from rgbd images. In *Computer Vision–ECCV 2012: 12th European Conference on Computer Vision, Florence, Italy, October 7-13, 2012, Proceedings, Part V 12*, pages 746–760. Springer, 2012. 2, 6, 7

[21] Julian Straub, Thomas Whelan, Lingni Ma, Yufan Chen, Erik Wijmans, Simon Green, Jakob J Engel, Raul Mur-Artal, Carl Ren, Shobhit Verma, et al. The replica dataset: A digital replica of indoor spaces. *arXiv preprint arXiv:1906.05797*, 2019. 2, 7

[22] Bin Tan, Nan Xue, Song Bai, Tianfu Wu, and Gui-Song Xia. PlaneTR: Structure-guided transformers for 3d plane recovery. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 4186–4195, 2021. 1, 2, 6

[23] Igor Vasiljevic, Nick Kolkin, Shanyi Zhang, Ruotian Luo, Haochen Wang, Falcon Z Dai, Andrea F Daniele, Mohammadreza Mostajabi, Steven Basart, Matthew R Walter, et al. Diode: A dense indoor and outdoor depth dataset. *arXiv preprint arXiv:1908.00463*, 2019. 2, 7

[24] Sagar M Waghmare, Kimberly Wilber, Dave Hawkey, Xuan Yang, Matthew Wilson, Stephanie Debats, Cattalyya Nuengsigkapian, Astuti Sharma, Lars Pandikow, Huisheng Wang, et al. Sanpo: A scene understanding, accessibility, navigation, pathfinding, obstacle avoidance dataset. *arXiv preprint arXiv:2309.12172*, 2023. 2, 7

[25] Karmesh Yadav, Ram Ramrakhya, Santhosh Kumar Ramakrishnan, Theo Gervet, John Turner, Aaron Gokaslan, Noah Maestre, Angel Xuan Chang, Dhruv Batra, Manolis Savva, et al. Habitat-matterport 3d semantics dataset. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 4927–4936, 2023. 2, 7

[26] Zehao Yu, Jia Zheng, Dongze Lian, Zihan Zhou, and Shenghua Gao. Single-image piece-wise planar 3d reconstruction via associative embedding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 1029–1037, 2019. 1, 2, 6

[27] Amir R Zamir, Alexander Sax, William Shen, Leonidas J Guibas, Jitendra Malik, and Silvio Savarese. Taskonomy: Disentangling task transfer learning. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 3712–3722, 2018. 2, 7