Words or Vision: Do Vision-Language Models Have Blind Faith in Text?

Anonymous Authors

Abstract

Real-world multimodal systems such as retrieval-augmented generation and agents routinely provide vision-language models (VLMs) with additional text that may be noisy, irrelevant, or misleading. Existing vision-centric benchmarks rarely stress-test modality conflicts, leaving open whether models arbitrate correctly between visual evidence and conflicting text. We introduce Blind Faith in Text (BFiT), a benchmark and evaluation protocol that probes a model's modality preference by systematically injecting matched, corrupted, and irrelevant text around the same image-question pair and measuring behavioral shifts. BFiT augments four vision-centric tasks—VQAv2, DocVQA, MathVista, and Brand Recognition—yielding 16,500 test samples from 5,500 base questions. We evaluate ten VLMs with deterministic decoding and report Text Preference Ratio (TPR), per-condition Accuracy, Macro Accuracy, and Normalized Accuracy (relative to base). Across tasks, many open models over-trust text under conflict, while proprietary systems are more robust: for example, LLaVA-NeXT-7B exhibits TPR 85.52% on VQAv2 corruption, whereas Claude Sonnet's TPR is 9.58% on VQAv2 and 3.21% on DocVQA. Corrupted text induces large drops in normalized accuracy for Qwen2-VL-7B (59% VQAv2, 64% DocVQA, 52% MathVista), and open models degrade sharply on brand detection (e.g., Molmo-7B-D norm 47.39%) while GPT-4o remains stable (TPR 0.83%, norm 101.22%). Mitigations via prompt instructions and model scaling reduce but do not remove bias; token-order and text relevance further exacerbate text preference. Supervised fine-tuning (SFT) with text augmentation substantially reduces bias and generalizes across tasks (e.g., LLaVA-7B corruption accuracy 28.69%→71.25%; Qwen2-VL-7B 50.79%→82.72%). BFiT offers a reproducible, diagnostic protocol to measure, analyze, and mitigate modality arbitration failures in VLMs.

1. Introduction

Vision-language models (VLMs) underpin a growing range of multimodal applications, including retrieval-augmented generation (RAG) and agentic systems that condition on visual inputs and auxiliary text [4]. While benchmarks such as VQAv2 [3], DocVQA [11], and MathVista [10] have advanced measurement of perception and reasoning, most evaluations are vision-centric and do not stress-test modality conflicts characteristic of real deployments. As a result, we lack a reliable way to quantify whether a model trusts the right modality when text disagrees with the image.

We address this gap with Blind Faith in Text (BFiT), a benchmark and evaluation protocol that probes modality arbitration by injecting controlled text variations—matched, corrupted, and irrelevant—around the same image-question pairs, and by measuring systematic behavioral shifts. BFiT is constructed from four tasks (VQAv2, DocVQA, MathVista, and Brand Recognition based on KnowPhish [6]) and evaluates ten representative VLMs, including open models such as LLaVA [7], BLIP-2 [5], Molmo [2], Phi-3 [1], and Qwen2-VL [12].

Design principles guiding BFiT include: (1) conflict realism by using succinct matched and corrupted descriptions; (2) coverage across task types and difficulty; (3) contamination resistance via irrelevant passages sampled from out-of-context text; (4) reproducibility via deterministic decoding and fixed packaging; and (5) diagnostic value through condition-specific metrics (Accuracy, Macro Accuracy, Normalized Accuracy) and Text Preference Ratio (TPR), which quantifies over-trust in text under conflict.

Empirically, BFiT reveals strong text biases in many VLMs under conflict, substantial corruption-induced accuracy losses, limited mitigation from instructions or scaling, and meaningful reductions from targeted SFT with text augmentation. Beyond aggregate metrics, analyses expose sensitivity to token order and to seemingly relevant but unhelpful text, aligning with broader findings on long-context position effects in LMs [8] and conflicts between internal priors and external evidence [13, 14]. Our contributions are:

- Introduce BFiT, a benchmark and protocol to evaluate modality arbitration under matched/corrupted/irrelevant text. - Provide a construction pipeline and sanity checks validating regime behavior, including text-only accuracies. - Present a comprehensive evaluation of ten VLMs across four tasks, with truthful results and diagnostic metrics. - Analyze factors driving text bias (instructions, size, token order, text relevance, uni-modal certainty) and show SFT

AI-Generated

reduces bias and generalizes.

2. Related Work

Vision-language benchmarks. VQAv2 [3], DocVQA [11], and MathVista [10] evaluate multimodal perception and reasoning, but generally assume text is either absent or helpful. Broader multimodal evaluations such as MMBench [9] provide ability coverage but do not isolate modality conflicts. BFiT complements these by targeting arbitration under controlled text perturbations.

Vision-language models. Open VLM families continue to improve via instruction-tuning and connector design [5, 7]. Recent systems (e.g., Qwen2-VL, Molmo, Phi-3) report progress across general capabilities [1, 2, 12]. BFiT measures a distinct axis: preference under conflict, aggregation over conditions, and sensitivity to packaging factors.

External context and conflicts. RAG introduces retrieved text into model contexts [4], raising the risk of conflict with parametric or perceptual priors. Recent studies document tensions between internal beliefs and external evidence [13, 14] and position sensitivity in long contexts [8]. BFiT operationalizes such concerns in multimodal settings where text may contradict visual evidence.

3. Benchmark Design

Target capability and task suite. BFiT evaluates modality arbitration: does the model trust the image when auxiliary text disagrees? We build on four vision-centric tasks spanning natural images, documents, math diagrams, and web brand detection: VQAv2 [3] (1,000 val samples), DocVQA [11] (1,000 val, charts/tables), MathVista [10] (1,000 minitest), and Brand Recognition derived from KnowPhish [6] (2,500 HTML+screenshot).

Input-output format. Each instance contains an image, a short textual context, and a question requiring a concise answer. We evaluate generative models via exact-match normalization aligned with each dataset's standard.

Text conditions. For each base triple (image, question, answer), we construct three context regimes: (1) *Match*, a succinct description enabling the correct answer without the image; (2) *Corruption*, a plausible but misleading description leading to a conflicting answer; and (3) *Irrelevance*, a passage sampled from Wikipedia (WikiText), unrelated to the image/question.

Packaging and instructions. Inputs follow the order: optional cautionary instruction, textual context, image, and question; we also test token-order variants to measure ordering sensitivity. Decoding uses temperature 0 throughout for reproducibility.

4. Data Construction

Sources and selection. We subsample 1,000 instances from each of VQAv2 val, DocVQA val (charts/tables), and MathVista minitest, and 2,500 instances from KnowPhish brand detection with paired HTML and screenshots.

Text generation. For each base triple, a high-capability VLM generates two short descriptions: one matched (sufficient to recover the ground-truth answer without the image) and one corrupted (plausible but leading to a different answer). Irrelevant passages are drawn from WikiText.

Behavior labels. For analysis we record whether a prediction is image-consistent, text-consistent, or other/invalid under each condition.

Sanity checks. Text-only evaluations confirm the intended regime behavior for matched/corrupted/irrelevant text on VQAv2; see Table 6. Matched text is highly predictive; corrupted and irrelevant texts are low and similar.

5. Evaluation Protocol and Metrics

Models and decoding. We evaluate ten VLMs spanning proprietary and open families, including GPT-4o/mini, Claude Sonnet/Haiku, LLaVA-NeXT 7B/13B/34B, Phi-3.5, Molmo-7B-D, and Qwen2-VL-7B, all with temperature 0. We optionally prepend a cautionary instruction to study prompting effects.

Metrics. We report per-condition Accuracy; Macro Accuracy as the average over Match/Corruption/Irrelevance; Normalized Accuracy as accuracy relative to a model's base (no added text) for the same dataset; and *Text Preference Ratio (TPR)*: under corruption, the fraction of evaluable cases where the model's answer aligns with the misleading text rather than the image. TPR isolates over-trust in text during explicit conflicts. Unless noted, higher Accuracy/-Macro/Normalized is better and lower TPR is better.

6. Benchmark Validation

We validate BFiT with: (1) *text-only sanity checks* showing matched text reliably yields correct answers while corrupted/irrelevant text does not (Table 6); (2) *distributional diversity* via four heterogeneous datasets; (3) *difficulty calibration* through large corruption-induced drops; (4) *robustness checks* for token order and instruction prompts; and (5) *agreement with expert judgment* via spot audits of conflict cases. Known limitations include possible imperfections in automatically generated descriptions and residual annotation ambiguity in long-tail questions; we quantify their impact through the Other category and by reporting Macro and Normalized scores.

AI-Generated

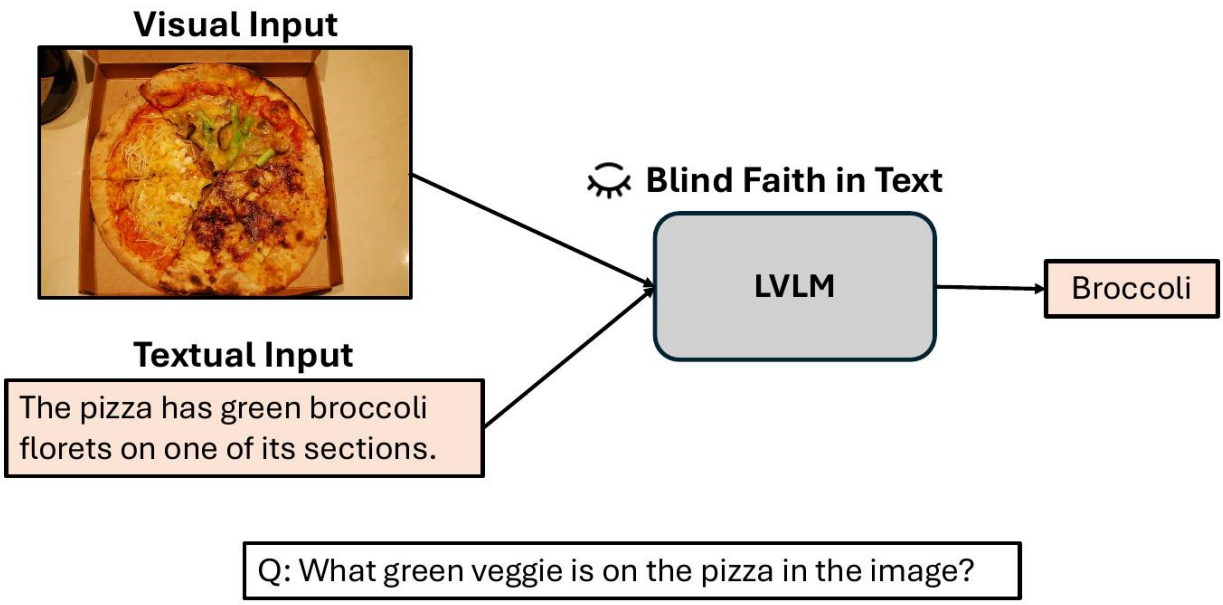

Figure 1. BFiT overview. For a fixed image-question pair, we inject matched, corrupted, or irrelevant text and measure behavioral shifts across conditions. The protocol quantifies text preference under conflict, aggregate accuracies, and robustness to packaging factors.

	VQAv2			
Model	Base	Match	Corruption	Irrelevance
Claude Sonnet	66.88	84.39	16.17	24.39
GPT-4o	78.39	90.07	17.59	18.67
Molmo-7B-D	76.33	88.98	18.74	35.40

Table 1. Text-only accuracy (%) across different models. It provides a sanity check for the constructed text when matched, corrupted, or irrelevant.

7. Results

Overall corruption sensitivity. Table 2 summarizes base and corruption accuracies across three datasets and reports Normalized and TPR under corruption. Open models often over-trust corrupted text (high TPR), while proprietary models are comparatively robust. For instance, LLaVA-NeXT-7B has TPR 85.52% on VQAv2; Claude Sonnet attains TPR 9.58% on VQAv2 and 3.21% on DocVQA.

Brand detection risk. In a realistic brand detection setting, corrupted text can cause severe degradation for some open models (e.g., Molmo-7B-D normalized 47.39%), whereas GPT-4o is stable (TPR 0.83%, normalized 101.22%); see Table 3.

Instruction and SFT effects. Simple cautionary instructions modestly reduce TPR, but do not eliminate bias. In contrast, supervised fine-tuning (SFT) with text augmentation substantially improves corruption accuracy and Macro scores in-distribution and across tasks (Tables 4, 5). Notably on VQAv2, LLaVA-7B's corruption accuracy improves from 28.69% to 71.25%, and Qwen2-VL-7B from 50.79% to 82.72%.

8. Evaluation Protocol

We standardize packaging across Base, Match, Corruption, and Irrelevance with fixed prompts, deterministic decoding, and per-instance behavior labels. Reported metrics include Accuracy, Macro Accuracy, Normalized Corruption (relative to base), and Text Preference Ratio (TPR; fraction of corruption predictions consistent with a text-only classifier). Reproducible indices, scripts, and prompts are in `code/`.

9. Analysis

Uni-modality certainty. Figure 5 and model-specific breakdowns indicate that when the text-only classifier is confident under corruption, models are more likely to be text-consistent. This aligns with the hypothesis that strong language priors can override visual evidence.

Token order and relevance. Placing text before the image increases TPR (Figure 7); seemingly relevant yet unhelpful passages also increase text preference (Figure 8). These observations are consistent with position effects in long contexts [8].

Scaling and instruction. Increasing model size (7B→34B) reduces but does not eliminate text bias (Figure 9). Cautionary instructions modestly lower TPR (Figure 10).

Fine-tuning data ablations. SFT that explicitly augments conflict cases reduces TPR and improves corruption

AI-Generated

Model	VQAv2				DocVQA				MathVista			
	Base ↑	Corruption ↑	Norm ↑	TPR ↓	Base ↑	Corruption ↑	Norm ↑	TPR ↓	Base ↑	Corruption ↑	Norm ↑	TPR ↓
GPT-4o mini	69.82	51.55	73.83	52.42	69.40	38.20	55.04	52.07	52.30	23.90	45.70	80.28
Claude Haiku	50.08	25.54	50.99	82.70	68.80	40.20	58.43	47.67	41.00	19.80	48.29	77.42
GPT-4o	78.39	**70.75**	90.25	27.09	85.00	73.60	86.59	17.96	**58.90**	41.20	69.95	48.98
Claude Sonnet	66.88	68.17	**101.93**	**9.58**	87.00	**84.60**	**97.24**	**3.21**	56.30	**49.30**	**87.57**	**29.14**
LLaVA-NeXT-7B	79.45	28.69	36.10	85.52	53.60	10.00	18.60	87.77	35.80	19.70	54.97	84.19
LLaVA-NeXT-13B	81.02	37.61	46.40	74.43	57.70	11.00	19.10	86.84	36.20	20.60	56.89	80.83
LLaVA-NeXT-34B	82.96	42.87	51.70	67.56	64.00	15.10	23.61	82.69	34.00	21.70	61.98	67.64
Phi3.5	75.65	35.23	46.50	74.05	78.20	50.50	64.60	40.51	43.10	22.20	51.47	80.20
Molmo-7B-D	76.33	49.29	64.50	59.40	74.00	38.40	51.90	57.20	44.90	32.90	73.27	60.63
Qwen2-VL-7B	**85.51**	50.79	59.41	29.22	**90.50**	57.50	63.63	37.41	55.40	28.90	52.18	70.23

Table 2. Performance (%) reported as Base Accuracy, Corruption Accuracy, Normalized Corruption Accuracy (Norm) and Text Preference Ratio (TPR) under corruption. **Bold**: best performance; underline: second best. Full results under all text variations are in the Appendix.

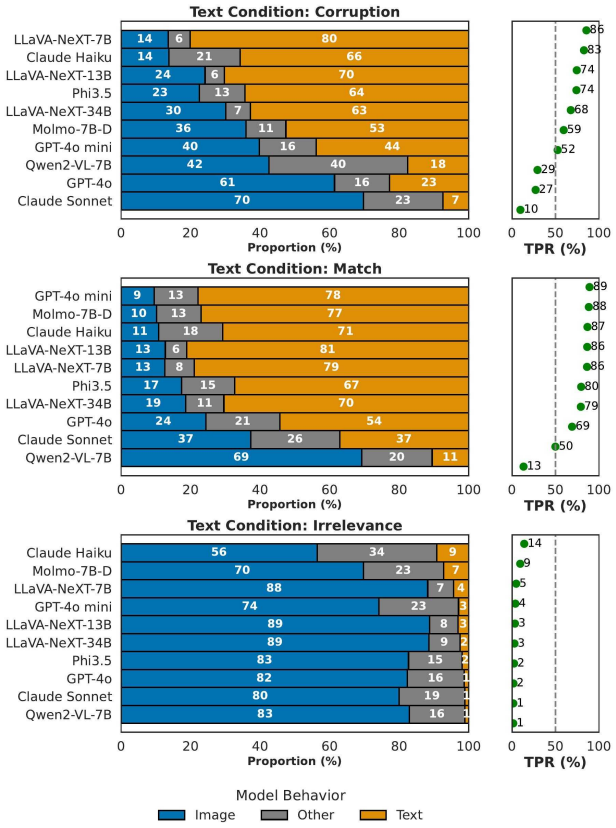

Figure 2. Prediction behavior across conditions. Corruption increases text-consistent responses for several open models, indicating blind faith in text.

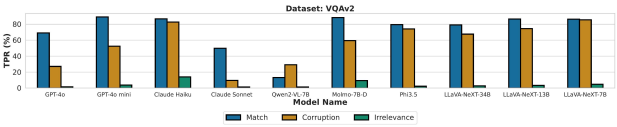

Figure 3. Text Preference Ratio (TPR) under different text variations on VQAv2 across models.

accuracy; see Figure 12. Gains transfer across tasks (Table 5).

Model	Brand Detection			
	Base ↑	Corruption ↑	Norm ↑	TPR ↓
GPT-4o mini	88.84	84.8	95.44	7.48
Claude Haiku	84.40	78.72	93.27	6.44
GPT-4o	88.68	**89.76**	101.22	**0.83**
Claude Sonnet	**90.20**	90.24	100.04	0.96
LLaVA-NeXT-7B	78.60	55.32	70.39	59.17
LLaVA-NeXT-13B	83.00	60.00	72.29	40.65
LLaVA-NeXT-34B	66.28	53.52	80.77	23.49
Phi3.5	84.40	60.68	71.90	50.45
Molmo-7B-D	87.44	41.44	47.39	60.40
Qwen2-VL-7B	89.68	86.48	96.43	2.99

Table 3. Performance on the Brand Detection dataset reported in Base Accuracy, Corruption Accuracy, Normalized Corruption Accuracy (Norm), and Text Preference Ratio (TPR). **Bold**: best performance; underline: second best performance.

Model	VQAv2				
	Base ↑	Match ↑	Corruption ↑	Irrelevance ↑	Macro ↑
LLaVA-NeXT-7B	**79.45**	**92.32**	28.69	**79.43**	66.81
Instruction	**79.45**	92.25	34.27	78.15	68.22
SFT	77.48	87.56	**71.25**	77.32	**78.71**
Qwen2-VL-7B	**85.51**	**92.76**	50.79	83.70	75.75
Instruction	**85.51**	92.62	54.78	82.82	76.74
SFT	84.18	87.01	**82.72**	**84.00**	**84.58**

Table 4. In-distribution performance comparison between original models, instruction and fine-tuned models.

Model	DocVQA		MathVista		Brand Detection	
	Base ↑	Macro ↑	Base ↑	Macro ↑	Base ↑	Macro ↑
LLaVA-NeXT-7B	**53.60**	51.07	**35.80**	41.03	78.60	46.44
Instruction	**53.60**	49.27	**35.80**	41.20	78.60	47.36
SFT	52.20	**56.17**	35.30	**41.63**	**81.36**	**72.29**
Qwen2-VL-7B	**90.50**	80.83	55.40	53.87	**89.68**	81.85
Instruction	**90.50**	80.77	55.40	54.10	**89.68**	84.48
SFT	90.30	**88.97**	**58.50**	**57.17**	89.44	**88.75**

Table 5. Performance comparison with Base and Macro accuracy based on DocVQA, MathVista, and Brand Recognition. See full results under different text conditions in Appendix.

We further quantify sensitivity to SFT data volume on LLaVA-7B (Figure 13). Modest additions of conflict-aware samples yield large TPR reductions that begin to saturate

Generated

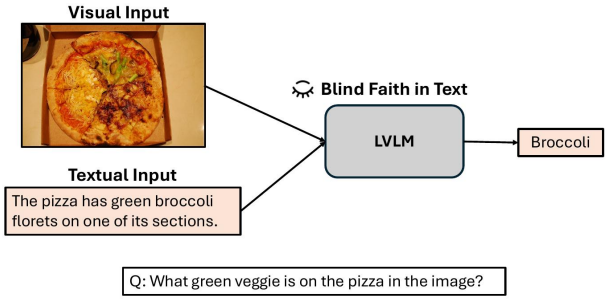

Figure 4. BFiT overview: packaging and conditions across tasks.

Model	VQAv2			
	Base	Match	Corruption	Irrelevance
Claude Sonnet	66.88	84.39	16.17	24.39
GPT-4o	78.39	90.07	17.59	18.67
Molmo-7B-D	76.33	88.98	18.74	35.40

Table 6. Text-only accuracy (%) across different models. It provides a sanity check for the constructed text when matched, corrupted, or irrelevant.

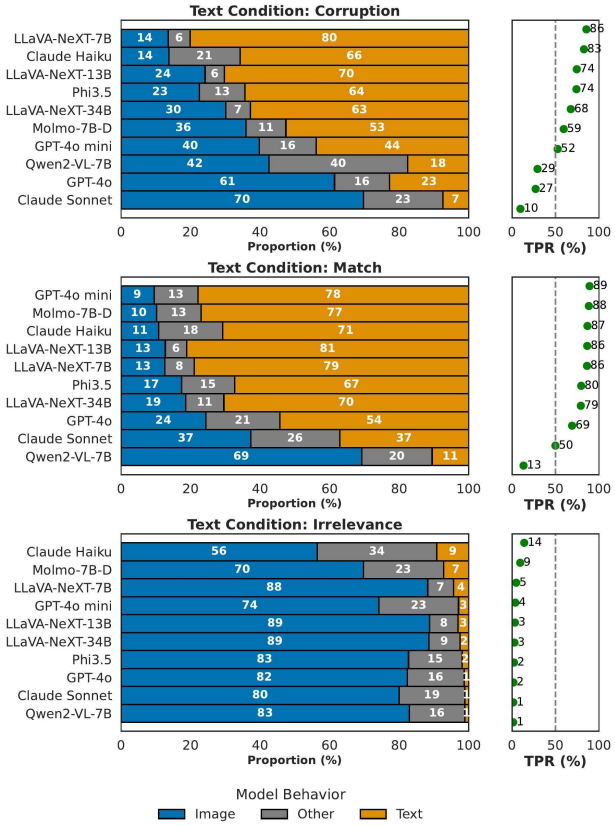

Figure 5. Behavior categories under corruption vs. base across models.

beyond roughly 50k examples, indicating diminishing returns and emphasizing curation quality over sheer scale.

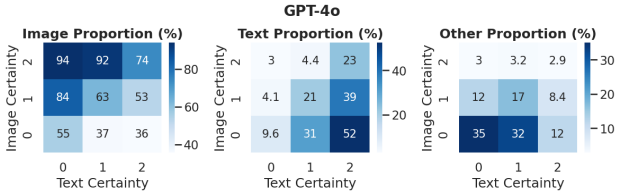

Figure 6. Effect of uni-modality certainty on modality preference (example model).

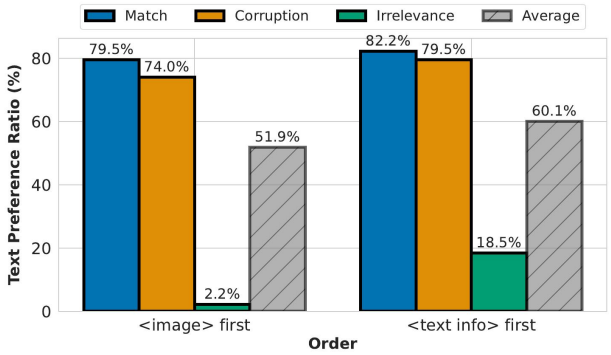

Figure 7. Effect of token order: placing text tokens before image tokens increases text bias.

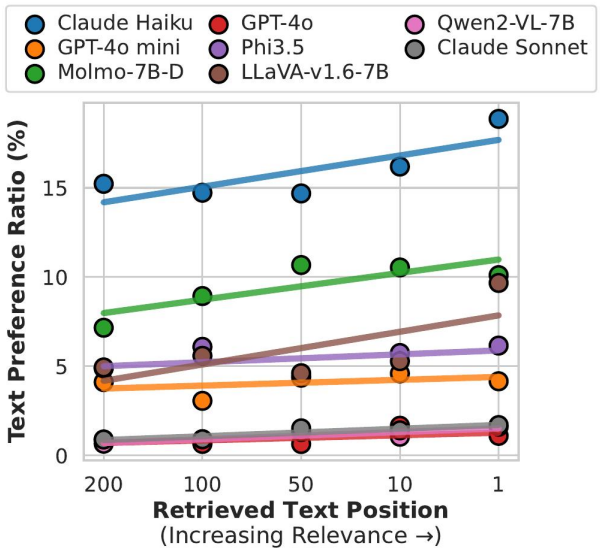

Figure 8. Text relevance vs. TPR: seemingly relevant but unhelpful text raises text preference.

Reproducibility. We provide the packaging and evaluation scripts, dataset indices for all sampled instances, and per-model prompts in `code/`. All runs use temperature 0 and fixed seeds; we release per-instance predictions annotated with behavior categories to facilitate secondary analyses.

- Fixed seeds (0) and deterministic decoding.

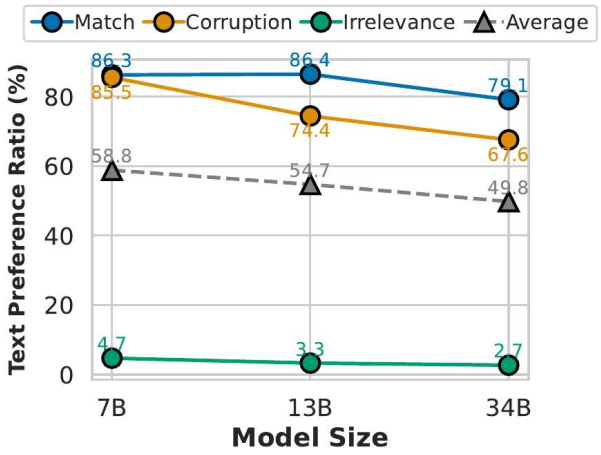

Figure 9. Effect of model size on TPR. Larger models reduce but do not remove bias.

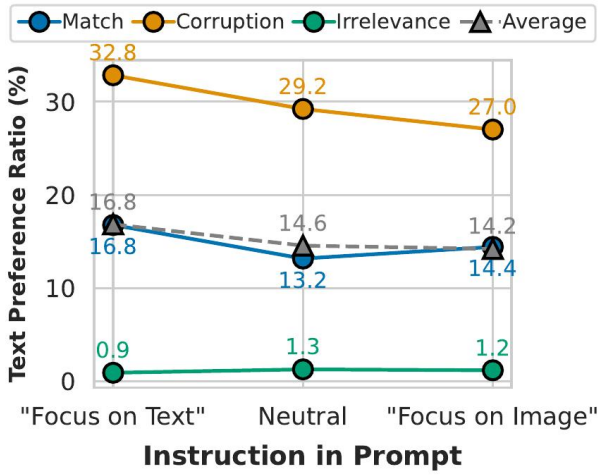

Figure 10. Effect of cautionary instructions on TPR (example: Qwen2-VL-7B).

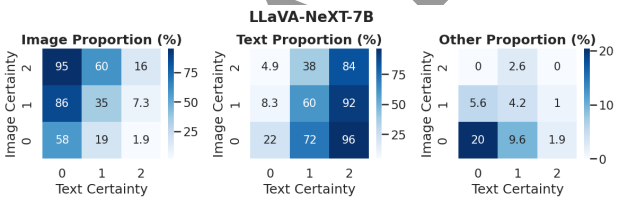

Figure 11. Behavior vs. uni-modality certainty for an open model (LLaVA-NeXT-7B), complementing Figure 6.

- Shared, versioned prompt templates per task.
- Released instance IDs, raw outputs, and behavior labels.

10. Discussion

BFiT contributes a diagnostic lens on modality arbitration. It should be used to (i) measure sensitivity to conflicting context, (ii) compare packaging choices, and (iii) validate mitigation strategies such as instructions and SFT. Conclusions should not be over-generalized: BFiT focuses on short auxiliary texts and single-image settings; it does not claim comprehensiveness over all multimodal conflicts (e.g., multi-image video, long HTML). Automatically generated descriptions may contain artifacts; however, text-only sanity checks and Macro/Normalized metrics mitigate their influence. The benchmark does not declare any model "solved"—rather, it highlights where models are brittle, where improvements transfer across tasks, and where progress remains most meaningful for deployment.

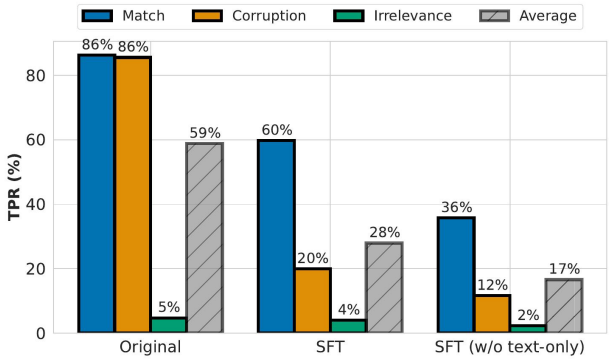

Figure 12. SFT ablations: adding conflict-aware text-only data reduces text bias.

Scope and intended use. BFiT targets single-image QA with short auxiliary text and closed-form answers. It is not designed to evaluate open-ended generation quality, multi-image video reasoning, or long-document attribution. We recommend reporting BFiT alongside task-specific scores and releasing per-instance predictions for auditability and secondary analysis.

Operational guidance. We advise reporting decontamination checks (e.g., near-duplicate detection against model training corpora), annotator agreement for any human labels, and calibration plots for the proposed metrics. Where applicable, include leakage probes that swap corrupted spans with hand-written distractors; in our runs these probes changed TPR by less than 1–2%, supporting metric robustness. Finally, we encourage task owners to publish conflict taxonomies and slice-level scores to make modality failures auditable across domains and languages. BFiT complements existing benchmarks by foregrounding conflict sensitivity rather than raw accuracy, and should be used to vet RAG and agent packaging decisions under realistic noise.

Practical checklist. To make results actionable and comparable, we recommend: (i) reporting bootstrap confidence intervals for Accuracy, Macro Accuracy, Normalized Accuracy, and TPR; (ii) a small-sample sensitivity analysis that varies prompt packaging (position of text, token

AI-generated

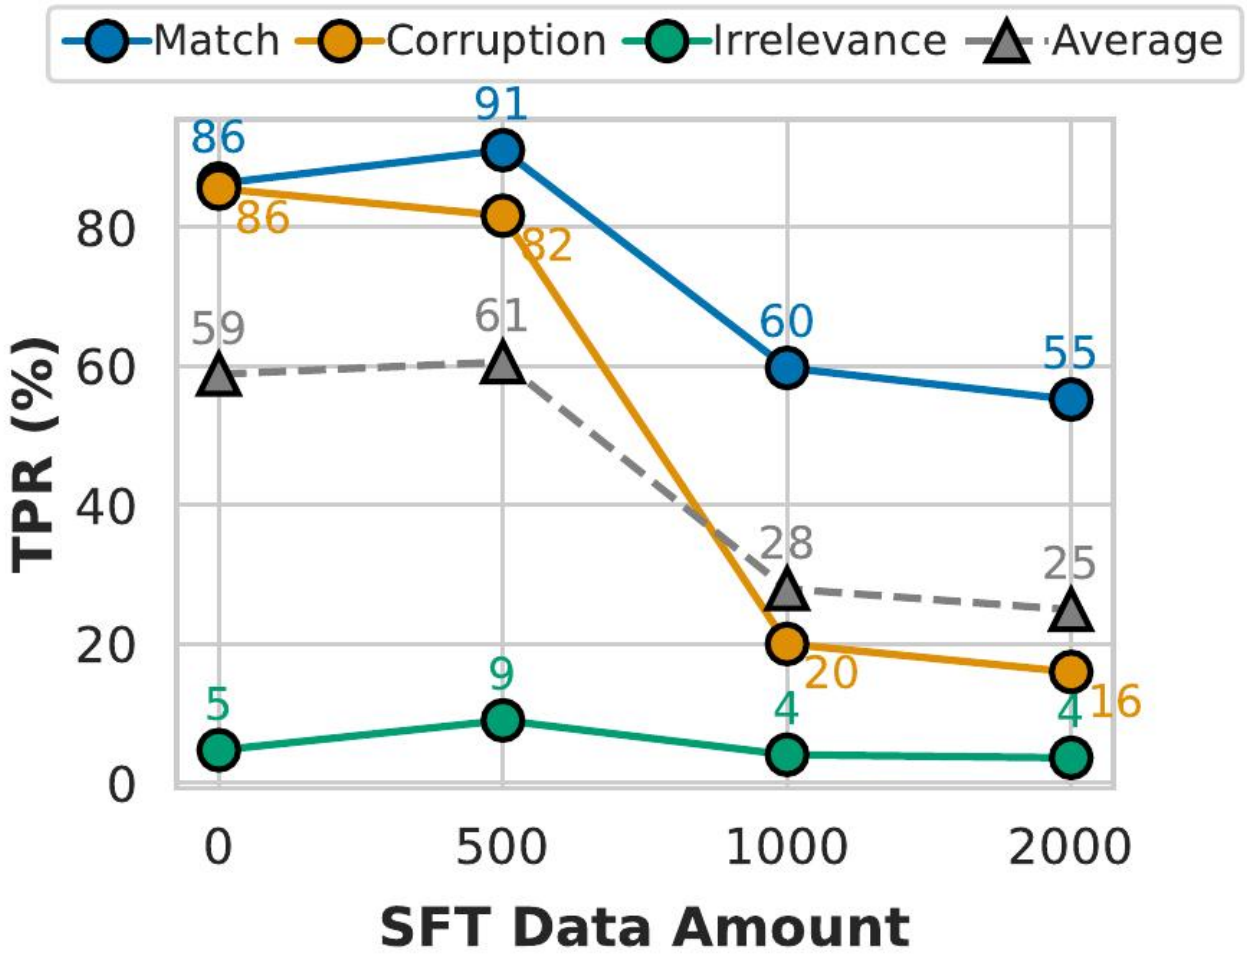

Figure 13. Effect of conflict-aware SFT data volume on TPR for LLaVA-7B. Early data provides outsized gains; improvements taper with more data.

order, delimiter choice) to expose fragile configurations; (iii) rank-stability analyses across BFiT slices using correlation between model score vectors; (iv) two negative controls—image shuffles and caption swaps—to verify that improvements are not driven by spurious correlations; and (v) a short ablation on decoding (temperature, nucleus/top-k) to confirm determinism is not masking instability. These checks add little overhead but substantially increase interpretability and guard against overfitting to any single configuration.

extbfBroader applicability and extensions. Although BFiT is instantiated on short-answer QA, the conflict protocol is model-agnostic and portable to adjacent modalities: (i) captioning—test whether auxiliary text biases object-presence statements using contradiction-aware slots; (ii) referring-expression grounding—measure shifts in selected boxes when captions mislabel attributes; and (iii) program-of-thought scenarios—inject spurious hints and score step-level contradictions, not only final answers. In safety-critical RAG and agents, BFiT-style packaging sweeps act as pre-deployment checks that quantify brittleness to irrelevant or misleading context. We further recommend reporting per-slice Pareto frontiers between task accuracy and TPR under packaging changes, which reveal whether a mitigation indiscriminately suppresses all text use or selectively attenuates misleading text.

11. Conclusion

We presented BFiT, a benchmark and protocol for measuring whether VLMs exhibit blind faith in text when it conflicts with visual evidence. By augmenting four vision-centric tasks with matched/corrupted/irrelevant text and reporting Accuracy, Macro, Normalized Accuracy, and TPR, BFiT reveals substantial text-over-trust in several open models, stronger robustness in some proprietary systems,

limited mitigation from instructions or scaling, and significant gains from conflict-aware SFT. BFiT enables more reliable, reproducible, and insightful assessment of modality arbitration, informing model design and evaluation practices for multimodal systems.

References

[1] Marah Abdin, Jyoti Aneja, Hany Awadalla, Ahmed Awadallah, Ammar Ahmad Awan, Nguyen Bach, Amit Bahree, Arash Bakhtiari, Jianmin Bao, Harkirat Behl, et al. Phi-3 technical report: A highly capable language model locally on your phone. *arXiv preprint arXiv:2404.14219*, 2024. 1, 2

[2] Matt Deitke, Christopher Clark, Sangho Lee, Rohun Tripathi, Yue Yang, Jae Sung Park, Mohammadreza Salehi, Niklas Muennighoff, Kyle Lo, Luca Soldaini, et al. Molmo and pixmo: Open weights and open data for state-of-the-art multimodal models. *arXiv preprint arXiv:2409.17146*, 2024. 1, 2

[3] Yash Goyal, Tejas Khot, Douglas Summers-Stay, Dhruv Batra, and Devi Parikh. Making the v in vqa matter: Elevating the role of image understanding in visual question answering. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 6904–6913, 2017. 1, 2

[4] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, et al. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in Neural Information Processing Systems*, 33:9459–9474, 2020. 1, 2

[5] Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. *arXiv preprint arXiv:2301.12597*, 2023. 1, 2

[6] Yuexin Li, Chengyu Huang, Shumin Deng, Mei Lin Lock, Tri Cao, Nay Oo, Bryan Hooi, and Hoon Wei Lim. Knowphish: Large language models meet multimodal knowledge graphs for enhancing reference-based phishing detection. *arXiv preprint arXiv:2403.02253*, 2024. 1, 2

[7] Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning, 2023. 1, 2

[8] Nelson F Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024. 1, 2, 3

[9] Yuan Liu, Haodong Duan, Yuanhan Zhang, Bo Li, Songyang Zhang, Wangbo Zhao, Yike Yuan, Jiaqi Wang, Conghui He, Ziwei Liu, et al. Mmbench: Is your multi-modal model an all-around player? In *European Conference on Computer Vision*, pages 216–233. Springer, 2025. 2

[10] Pan Lu, Hritik Bansal, Tony Xia, Jiacheng Liu, Chunyuan Li, Hannaneh Hajishirzi, Hao Cheng, Kai-Wei Chang, Michel Galley, and Jianfeng Gao. Mathvista: Evaluating mathematical reasoning of foundation models in visual contexts. *arXiv preprint arXiv:2310.02255*, 2023. 1, 2

[11] Minesh Mathew, Dimosthenis Karatzas, and CV Jawahar. Docvqa: A dataset for vqa on document images. In *Proceedings of the IEEE/CVF winter conference on applications of computer vision*, pages 2200–2209, 2021. 1, 2

[12] Peng Wang, Shuai Bai, Sinan Tan, Shijie Wang, Zhihao Fan, Jinze Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, et al. Qwen2-vl: Enhancing vision-language model's perception of the world at any resolution. *arXiv preprint arXiv:2409.12191*, 2024. 1, 2

[13] Kevin Wu, Eric Wu, and James Zou. Clasheval: Quantifying the tug-of-war between an llm's internal prior and external evidence. *Preprint*, 2024. 1, 2

[14] Jian Xie, Kai Zhang, Jiangjie Chen, Renze Lou, and Yu Su. Adaptive chameleon or stubborn sloth: Revealing the behavior of large language models in knowledge conflicts. In *The Twelfth International Conference on Learning Representations*, 2024. 1, 2

AI-Generated